

Mitigating Adverse Selection in Concentrated Liquidity AMMs with Dynamic Fees: An Agent-Based Model Approach

Daniele Maria Di Nosse^{†1} and Fabrizio Lillo¹

¹Scuola Normale Superiore, Pisa, Italy

June 11, 2026

[†]Corresponding author. Email: daniele.dinosse@sns.it

Abstract

Automated Market Makers (AMMs) based on concentrated liquidity, such as Uniswap v3, significantly improve capital efficiency but expose Liquidity Providers (LPs) to adverse selection costs, formalized as Loss-Versus-Rebalancing (LVR). While theoretical literature quantifies these costs, the interplay between realistic blockchain microstructure and endogenous pricing mechanisms remains under-explored. This paper develops a granular Agent-Based Model (ABM) of a Uniswap v3 pool interacting with a stochastic reference market governed by Heston volatility dynamics. The framework incorporates discrete block propagation, mempool latency, and a heterogeneous population of agents, including latency-sensitive arbitrageurs, smart routers, Maximal Extractable Value (MEV) searchers, and active LPs benchmarked against a frictionless rebalancing strategy. We propose and evaluate dynamic fee schedules driven by volatility and order-flow toxicity proxies intended to internalize the cost of adverse selection. Our simulations investigate the conditions under which LPs can achieve positive hedged PnL (fees minus LVR). The analysis suggests that dynamic fee adjustments can successfully mitigate arbitrage losses during high-volatility regimes, indicating that adaptive pricing is essential for the long-term economic sustainability of passive liquidity provision.

Contents

1	Introduction	2
2	Literature Review	4
2.1	Analytical and Theoretical Approaches	5
2.2	Empirical Studies	5
2.3	Agent-Based Modeling and Simulations	6
2.4	ABMs for Market Design and Financial-Policy Analysis	8
2.5	Dynamic fees in deployed AMMs	8
3	Model	9
3.1	Uniswap v3 Market Mechanism	9
3.1.1	Discrete price grid, ticks, and square-root price	9
3.1.2	Liquidity positions and active liquidity	10
3.1.3	Swap dynamics within a tick	10

3.1.4	Tick crossing and pathwise execution of large swaps	11
3.1.5	Mempool	12
3.2	Reference Market (CEX)	12
3.2.1	Heston Volatility Model	13
3.2.2	Market Impact	13
3.3	LVR and PnL Accounting	13
3.3.1	Rebalancing Benchmark in Discrete Time	15
3.4	Agents	15
3.4.1	Arbitrageurs	16
3.4.2	Liquidity Providers	17
3.4.3	Smart router	21
3.4.4	Noise Trader	22
3.5	Dynamic Fee Schedules	22
3.6	Initialization	24
4	Results	24
4.1	Microstructure diagnostics	27
4.2	Cross-scenario outcome summary	29
4.3	LVR vs fees vs block size	32
5	Conclusions	37

1 Introduction

Automated market makers (AMMs) have emerged as a key innovation in Decentralized Finance (DeFi), allowing users to trade assets through liquidity pools instead of traditional order books. From a market-microstructure perspective, AMMs can be interpreted as automated dealer markets in which the fee plays a role analogous to a spread or transaction cost, while inventory adjustment is embedded in the pricing rule rather than in discretionary dealer quotes [25, 6, 5, 10]. In AMMs like Uniswap, traders swap against pooled assets and pay fees to liquidity providers (LPs) who supply those assets. Uniswap v3 introduced concentrated liquidity, enabling LPs to allocate capital within specified price ranges to improve fee earnings and capital efficiency [2]. This advancement comes with increased exposure to market risk: if the price moves outside an LP’s range, the position can become effectively inactive and fully converted to one asset. Thus, while concentrated liquidity boosts potential fee revenue, it also exacerbates the impermanent loss risk that can be expressed as the opportunity cost or loss an LP incurs compared to simply holding the assets. A central question is whether the fees earned by providing liquidity are sufficient to compensate for these risks, i.e. whether liquidity provision is profitable net of hedging costs.

Recent research [32] has formalized the notion of LP profitability by introducing the concept of “hedged PnL”, which isolates the profit-and-loss of an LP from the market risk. In particular, Milionis et al. (2023) define the delta-hedged LP return as the difference between the so-called *loss-versus-rebalancing* (LVR) cost and the trading fees earned [32]. LVR is the loss an LP suffers due to adverse selection: because AMMs execute trades at stale prices, informed arbitrageurs extract value when prices move, leaving LPs worse off than if they had continuously rebalanced their portfolio at market prices. This adverse-selection interpretation parallels the classical market-making literature, where liquidity suppliers require compensation for inventory risk and informed order flow through spreads and execution premia [25, 6, 12, 21]. In effect, LVR measures the gap between the LP’s value and a benchmark rebalancing strategy’s value (analogous to a mark-to-market loss).

The hedged PnL, given by fees earned minus LVR, represents the LP’s real economic profit after accounting for the cost of maintaining a market-neutral position (i.e. after hedging out pure price exposure) [32]. This metric has become a preferred indicator of true profitability because an LP’s unhedged returns are largely driven by overall asset price movements (which could be obtained without providing liquidity by simply holding the assets). By contrast, a delta-hedged LP’s returns reflect only the value added (or lost) by liquidity provision, separating out market risk.

The profitability of liquidity provision in concentrated AMMs is not merely a theoretical concern. It has significant practical implications for traders, investors, and protocol designers. Empirical evidence has raised concerns that, for many volatile asset pools, LPs on Uniswap v3 underperform a buy and hold strategy. In other words, the combination of impermanent loss and LVR can outweigh the fees earned. For example, an in-depth analysis of Uniswap v3 by Heimbach et al. [23] found that across major volatile pools (like ETH-USD or BTC-ETH pairs), LPs on average earned negative returns relative to holding, especially when they provided very narrow (concentrated) ranges. The mean daily returns of LP positions in such pools were below zero, indicating that fees did not fully offset the losses from adverse price movements and arbitrage. Only in low-volatility stablecoin-stablecoin pools did LPs consistently earn small positive returns, as impermanent loss was negligible in those cases. Likewise, a comprehensive empirical study by Fritsch and Canidio [19] examined many large pools and concluded that losses to arbitrageurs (LVR) exceeded fee revenues in most cases, resulting in negative hedged PnL for LPs on those pools. [The empirical evidence is therefore consistent with the theoretical view that AMM fees must compensate not only for price volatility, but also for the timing and informational content of order flow \[32, 19, 28\].](#) Notably, they report that Uniswap v3’s concentrated liquidity pools were less profitable for passive LPs than the earlier Uniswap v2 pools, despite the higher fee potential, due to the greater arbitrage losses in v3. These findings underscore that liquidity provision, particularly in volatile markets, can be an adverse trade for uninformed or passive providers. It also highlights the importance of active management or hedging strategies to improve outcomes.

Against this backdrop, understanding the drivers of LP profitability and exploring methods to enhance it has become a multidisciplinary endeavor. On one hand, analytical models from finance can quantify the expected costs and returns of liquidity provision under various conditions. On the other hand, Agent-Based Models (ABMs) and simulations offer a way to study complex interactions in AMM ecosystems between various kinds of agents that are difficult to capture with closed-form theory. The benefit of ABMs is their description of complex systems from the bottom up by specifying the behaviour (often very simple) of individual agents and letting aggregate properties emerge from their interactions. The main drawback, however, lies in the large number of parameters on which these models typically rely, whose calibration is often nontrivial. Broadly speaking, ABM can be interpreted in two distinct ways. They can be viewed either as data-generating mechanisms designed to reproduce as many stylized facts as possible (for example, as synthetic market generators useful for testing trading strategies) or they can serve as controlled and realistic laboratories in which the researcher studies the qualitative response of the system to variations in specific parameters (for instance, how a change in tick size affects the price dynamics of a financial asset). [This use of ABMs as controlled laboratories is well established in artificial financial-market research, where heterogeneous-agent simulations have been used to study transaction taxes, market microstructure, high-frequency trading, cancellation rules, fragmentation, and the propagation of external information \[36, 13, 39, 30, 35, 26, 9, 33\].](#) In this paper, we adopt the latter perspective. Our goal is to investigate whether, and in what way, dynamically adjusted fee schedules influence the economics of liquidity provision in a Uniswap v3-like AMMs. [In this sense, AMM fee adaptation is the DeFi analogue of a market-design intervention: changing a transaction cost or spread-like rule and observing how liquidity, routing, arbitrage, and welfare-relevant outcomes respond in a](#)

heterogeneous-agent system [39, 30, 35, 33].

This paper presents an ABM for a Uniswap v3-style pool that allows agents to interact each other, making decisions based on the state of the pool and a reference market, a Centralized Exchange (CEX). Specifically, the model incorporates:

- tick-based concentrated liquidity;
- a reference CEX price evolving according to a Heston volatility dynamics with permanent price impact due to arbitrageurs;
- a mempool-based execution;
- heterogeneous agents: arbitrageurs, noise traders, liquidity providers and Maximal Extractable Value (MEV) searchers;
- passive and active LPs with budgets, review clocks, and risk management;
- dynamic taker fees driven by volatility or order flow toxicity;
- a rebalancing benchmark for each LP that enforces wealth conservation and allow for a clean hedged-vs-unhedged PnL decomposition;

The primary objective is to understand how (and if) dynamical liquidity taker fees can be beneficial for liquidity provision, reducing LVR and/or increasing fee revenue for the LPs.

These findings extend the analytical and empirical AMM literature on predictable loss, LVR, and arbitrage losses by studying how adaptive fee rules interact with block-time execution and strategic liquidity provision [32, 11, 10, 19]. The simulations deliver three main findings. First, under block-time execution, static fees are insufficient to compensate standing LPs for stale-quote adverse selection, especially when liquidity is concentrated. Second, dynamic fees can improve hedged LP profitability, but their effectiveness depends on the information used to adjust the fee: rules that react to the contemporaneous gap between the AMM price and the external reference price perform better in our experiments than rules based only on recent price volatility. Third, fee adaptivity creates an endogenous trade-off: it protects standing liquidity by widening the fee wedge, but it also reduces smart-router DEX share and, when very short-lived liquidity provision is allowed, can redirect part of the enlarged fee pool toward latency-sensitive liquidity providers.

The paper is organized as follows. Section 2 provides a brief literature review on liquidity provision in AMMs, with particular attention to LP profitability, adverse selection, and dynamic fee mechanisms. Section 3 presents the agent-based model, describing the Uniswap v3 market mechanism, the reference market, the LP PnL accounting framework, the agent population, and the dynamic fee schedules considered in the analysis. Section 4 reports the main simulation results, including the microstructure diagnostics and the comparison across fee policies and model specifications. Finally, Section 5 concludes and discusses possible directions for future research.

2 Literature Review

This section reviews the main strands of literature relevant to our research question. We first discuss analytical contributions that formalize LP profitability and adverse selection in AMMs, then summarize the empirical evidence on realized LP performance, and then examine simulation-based and agent-based approaches for AMMs. We then add a broader perspective on ABMs as tools for market-design and policy analysis [36, 13, 33]. We conclude with a brief overview of academic and deployed dynamic fee mechanisms, since these designs provide useful guidance for the fee rules analyzed in our model.

2.1 Analytical and Theoretical Approaches

The foundational models of AMMs and liquidity provider returns have benefited from insights in financial economics. A seminal theoretical contribution is the Black-Scholes-style model for AMM liquidity proposed by Milionis et al. (2023), which explicitly quantifies LP returns and the LVR cost in continuous time [32]. By modeling the AMM as replicating a continuously rebalanced portfolio, they derive closed-form expressions for LVR under diffusion assumptions. Their model shows that an LP’s position can be viewed analogously to writing options: the LP continuously “sells low and buys high” when prices move, akin to being short volatility. This explains why, even after hedging out directional price risk, LPs face a systematic shortfall (LVR). The analytical result $extHedgedPnL = extFees - extLVR$ summarizes how fees must overcome LVR for an LP to profit [32].

This framing connects AMM liquidity provision to the broader market-making literature, where spreads and execution policies compensate dealers for inventory risk, asymmetric information, and adverse selection [25, 6, 12, 21]. Cartea, Drissi, and Monga study AMM execution, speculation, predictable loss, and optimal liquidity provision, providing a direct bridge between DeFi mechanisms and continuous-time market-making models [10, 11]. Their work ties into the idea of setting optimal trading fees to balance trade frequency and arbitrage loss: higher fees protect LPs from being picked off by arbitrage (reducing LVR), but too high fees reduce volume and fee income. More generally, equilibrium and stochastic-control formulations reinforce the same trade-off: without an appropriate fee or execution design, liquidity-provider returns tend to be squeezed by informed trading and arbitrage in volatile markets.

A distinct line of theoretical inquiry addresses strategies for mitigating impermanent loss. Some works explore hedging impermanent loss using financial derivatives or alternative mechanisms. For example, Maire and Wunsch [28] propose a market-neutral liquidity provision strategy where the LP simultaneously holds an appropriate portfolio of futures/options to offset AMM exposure. While this approach can eliminate directional risk and LVR, they find it requires substantial extra capital and its efficacy depends on market conditions. Another idea is to redesign the AMM mechanism itself; Canidio and Fritsch examine batched trading/auctions as a protocol design to reduce LVR and sandwich attack risks, thereby improving LP outcomes [8]. These analytical studies contribute to a broader understanding that LP profitability is fundamentally tied to market microstructure – the timing of price updates, arbitrage speed, fee levels, and volatility all determine whether fees outweigh the inherent losses from providing liquidity.

2.2 Empirical Studies

Empirical research has leveraged blockchain data to measure realized LP performance in AMMs. Using transaction and pool data from Uniswap and other platforms, researchers can compute the actual returns of LP positions over time and compare them to benchmarks. The aforementioned study by Heimbach et al. [23] is one of the first comprehensive looks at Uniswap v3 LP data. They tracked individual LP positions in various pools and calculated their holding-period returns, inclusive of fees and impermanent loss, versus just holding assets. Their findings were striking: for volatile pairs like WETH/USDC or WBTC/WETH, a majority of LPs would have been better off not providing liquidity. Narrow-range positions (which concentrate liquidity aggressively) exhibited higher fee income but also much higher impermanent losses, often leading to negative net returns on average. By contrast, in stablecoin pools (e.g. DAI/USDC), volatility is so low that impermanent loss was minimal and nearly all LPs earned small positive returns (albeit with lower fee yields). Heimbach et al. also observed that LP positions with very short lifetimes saw the most extreme

outcomes – some made high returns by timing volatility/volume spikes correctly, while many others lost significantly – suggesting that active management and informed timing are crucial for success in Uniswap v3’s environment.

Subsequent empirical studies have reinforced these conclusions. Fritsch and Canidio empirically quantify the total dollar value that flowed from LPs to arbitrageurs due to LVR across many Uniswap pools [19]. They find that this arbitrage extraction often exceeded the fees paid to LPs, resulting in net losses. Their work also highlights chain speed as a factor: on faster blockchains (or hypothetical faster Ethereum), the arbitrage gaps – and hence LVR – would shrink, improving LP outcomes (since prices update closer to continuously). Another empirical analysis by Heimbach and Wattenhofer (2023) looked into how fee tier selection and volatility interact, noting that LPs choosing higher fee tiers (e.g. 1% fee) experienced smaller impermanent losses but at the cost of lower trade volume (fewer fees), whereas low fee tiers attracted more volume but inflicted more IL; neither extreme universally dominated, implying an optimal fee tier might depend on the volatility and trading intensity of the pair. Overall, empirical evidence paints a cautious picture: passive liquidity provision is often a negative expected value proposition for volatile asset pairs, whereas for more stable pairs or with very active management (frequently rebalancing ranges), LPs can achieve modest gains.

2.3 Agent-Based Modeling and Simulations

To capture the complex dynamics of AMM ecosystems, researchers have increasingly turned to agent-based models (ABM) and simulations. In an ABM, one can simulate a population of agents with different roles – e.g. market-makers (LPs), traders, arbitrageurs – interacting according to set rules. This approach, common in econophysics, allows experimentation with various scenarios (different volatility regimes, fee levels, liquidity distributions, agent strategies, etc.) and observation of emergent outcomes like LP profitability or market stability.

One early example of applying simulations to AMMs is by Angeris et al. [5], who not only provided a theoretical analysis of Uniswap but also created an agent-based simulation to test market conditions and arbitrage behavior in a Uniswap pool. Their simulation (using the Gauntlet platform) included agents trading and arbitrageurs ensuring prices don’t deviate from an external reference. The results supported theoretical properties (e.g. Uniswap’s pricing converges to market price via arbitrage) and hinted at how LP returns depend on volatility and arbitrage frequency, laying groundwork for later studies.

A more focused ABM study on LP profitability is by Cohen et al. [14]. They first derive a theoretical upper bound on fees that would be required for an LP to fully hedge away market risk in a constant-function market maker (CFMM) like Uniswap. They show analytically that, under typical conditions, the required fee to break even (ignoring risk premia) is quite high, hinting that standard fee levels (e.g. 0.3%) might be insufficient. They then implement a multi-agent simulation with LPs and traders, varying parameters such as the fee level, asset volatility, and the rate of trade (liquidity taker arrival). The simulation confirms that on average, fee income is insufficient to cover the hedging cost for LPs, especially as volatility increases or trade frequency decreases [14]. In their experiments, even raising fees does not fully fix the problem: while higher fees reduce LVR (by dissuading marginal arbitrage trades), they also reduce volume. The ABM provides a nuanced view of how different factors trade off, illustrating, for example, that at very high volatility, no reasonable fee can protect LPs, whereas at low volatility, even a small fee can yield positive expected hedged PnL. This aligns with the empirical notion that LP profitability is highly regime-dependent.

Another agent-based approach is seen in Hafner and Dietl [22]. They conduct simulations to determine under what conditions an LP’s position outperforms a holding strategy. Their model

iterates over one-year periods with agents simulating trades consistent with historical Uniswap volumes and volatility. Intriguingly, their findings suggest that if the asset’s price fluctuations are not too extreme (for example, not more than a -75% drop or +300% increase over the year) and the fee volume is substantial, then an LP could realize a net profit relative to simply holding the assets [22]. Essentially, in moderate market conditions with decent trading activity, fee accumulation can outpace impermanent loss. However, when markets are highly volatile (large price divergence) without commensurate volume, the LP is likely to underperform holding. This result offers a more optimistic counterpoint, indicating that there are realistic scenarios (especially for less volatile assets or shorter time horizons) where liquidity provision pays off. It underscores the importance of market environment in any evaluation of LP profitability.

Agent-based models have also been used to design and test strategic LP behaviors. Rather than treating LPs as passive liquidity suppliers, some works model them as active agents who can adjust their liquidity ranges or withdraw/react based on market signals. Fan et al. [18] present a game-theoretic and computational approach for dynamic liquidity provision in Uniswap v3. They consider a setting with one strategic LP who periodically updates their price range to maximize earnings, anticipating price moves. By simulating price paths and using techniques like reinforcement learning or neural-network-based optimization, they show that active strategies (like regularly resetting liquidity ranges based on volatility or price trends) can significantly improve an LP’s PnL compared to a static strategy. Their work, along with others, suggests that liquidity provision in Uniswap v3 is an active trading strategy akin to market-making in traditional markets – requiring optimization and sometimes rapid response to market changes. This blurs the line between LPs and traders: the most successful LPs behave like informed agents, adjusting positions to manage inventory and risks.

Incorporating hedging into agent strategies is another frontier. Zhang et al. [41] employ deep reinforcement learning (DRL) to train an agent that provides liquidity in Uniswap v3 while simultaneously hedging with external instruments (like futures) to offset price risk. The DRL agent learns to shift its liquidity range and maintain an appropriate hedge ratio in a way that maximizes the expected hedged PnL (trading off fees earned against hedging costs and gas fees). In simulations on volatile pairs like ETH/USDC, the learned strategy outperformed naive strategies and achieved better risk-adjusted returns [41]. This approach illustrates the potential of ABM combined with machine learning: one can simulate a realistic market environment (with price paths, trading costs, etc.), train agent strategies within it, and derive insights on best practices for LPs. Notably, Zhang et al. confirm that hedging significantly reduces PnL volatility (by removing most impermanent loss), and the remaining key to profitability is optimizing fee collection versus transaction costs. Their agent effectively implements the delta-hedged LP strategy in an automated fashion, reinforcing the concept that hedged PnL is the right objective for liquidity providers.

In summary, the literature on liquidity provision profitability spans analytical finance, empirical data analysis, and agent-based simulation. The consensus from these varied approaches is that liquidity provision in concentrated AMMs is a high-risk, low-margin endeavor for most typical users, especially in volatile markets. Adverse selection (LVR) tends to systematically eat into LP returns, meaning that without either specialized strategies or favorable market conditions, LPs may not be adequately compensated for the risk. Agent-based models have been particularly useful in stressing the system under different conditions and testing potential improvements. They have shown, for instance, that protocol design changes (like fee adjustments or batch auctions) can alter the balance of profits, and that sophisticated, active strategies (possibly aided by algorithmic agents or automation) are necessary to achieve positive expected returns. From a finance perspective, this aligns with the idea that passive market-making can be a losing proposition when trading with informed counterparts – an insight that echoes traditional market microstructure theory (where market makers require a spread or other edge to profit). In the context of AMMs, the “spread” is

essentially the fee, and current fee levels often appear too low to offset the adverse selection cost under fast-moving prices.

2.4 ABMs for Market Design and Financial-Policy Analysis

A related strand of artificial financial-market research uses ABMs to evaluate market-design interventions. Early artificial markets showed how stylized price dynamics can emerge from heterogeneous interacting agents [36, 13]. Subsequent work uses ABMs to study how transaction taxes and fee-like frictions affect volatility, liquidity, and market stability under different microstructures [39, 30, 35]. Other studies examine how rule changes related to hedging instruments, high-frequency trading, order cancellation, latency, and external information can amplify or dampen market instability [7, 26, 9, 20, 40]. Our use of an ABM to test dynamic AMM fees follows this market-design logic: the fee rule is treated as a policy lever whose impact depends on heterogeneous agent responses and feedback through prices, routing, liquidity provision, and arbitrage [33].

2.5 Dynamic fees in deployed AMMs

Academic work on AMM fee choice emphasizes that the fee is not merely a revenue parameter but a control variable that trades off volume against adverse selection and arbitrage losses [32, 11]. Deployed protocols implement this idea through simpler, bounded controllers based on imbalance, volatility, or activity proxies.

Beyond academic proposals, several large AMM families already deploy fee schedules that adapt endogenously to risk or imbalance conditions. These implementations are useful reference points because they reveal which signals are considered sufficiently robust to run “on-chain” and which functional forms are preferred to keep fees bounded, smooth, and hard to game.

A first design pattern is off-peg / imbalance-based fees, common in stable-asset designs. The basic idea is that, for assets that should trade close to a target relative price, the pool’s reserve composition is itself an informative risk signal. When the pool is close to its target balance, trades are likely to be ordinary liquidity-taking flow and can be charged a low baseline fee. When trades push the pool away from balance—for example, leaving it with too much of one asset and too little of the other—the pool is effectively absorbing inventory and de-peg risk. The fee is therefore increased to discourage trades that further worsen the imbalance and to compensate LPs for taking the other side of stressed flow. In Curve’s Stableswap-NG, this idea is implemented with a compact balance metric $K(x, y) = \frac{4xy}{(x+y)^2} \in (0, 1]$, where x and y denote normalized reserves and $K = 1$ at perfect balance ($x = y$). As the pool becomes more imbalanced, K falls below one and the dynamic fee rises through a rational function that equals the baseline fee at the peg and approaches a capped multiple of it as imbalance grows [16]. Curve v2 (CryptoSwap) implements a closely related idea: the fee interpolates between a “mid” fee (near balance) and an “out-of-range” fee (far from balance) through a smooth imbalance coefficient, often expressed via a geometric-vs-arithmetic-mean ratio and a reduction factor that shrinks under imbalance [15]. For modeling purposes, both mechanisms can be read as: measure how far the pool state is from its target composition with a scalar statistic, then map that statistic monotonically into a higher effective fee, with a hard ceiling.

A second pattern is activity/volatility-accumulator fees, prevalent in “bin” AMMs. Trader Joe’s Liquidity Book computes fees as the sum of a base term proportional to the bin step and a quadratic variable term driven by a volatility accumulator that grows when swaps traverse bins and decays when activity slows [27]. Meteora’s DLMM follows the same blueprint (base plus a capped quadratic variable component) with explicit scaling constants and a maximum fee rate [31]. Conceptually,

these designs make fees respond quickly to short-lived bursts of toxicity (rapid price moves and frequent bin crossings), while ensuring mean reversion through time-decay.

A third pattern appears in concentrated-liquidity designs that remain close to the Uniswap v3 paradigm. Algebra’s adaptive fee combines (i) a volatility premium computed as a sum of sigmoid responses of a tick-based volatility estimator and (ii) a “volume-per-liquidity” regulator, then clamps the result between minimum and maximum fees [3, 4]. The practical takeaway is that smooth saturating functions (sigmoids) are used to avoid fee flicker, while multiplicative regulators incorporate local market conditions (how much flow is being processed per unit of liquidity).

Overall, deployed protocols converge on three robust implementation choices: (i) a low-dimensional risk signal (imbalance, volatility proxy, bin-crossing intensity), (ii) a monotone mapping that increases fees in stressed regimes (often convex/quadratic or saturating), and (iii) explicit caps and decay to stabilize dynamics. These motifs directly motivate the simplified controllers we test in Section 3.5, where fees are updated on a block clock using volatility- and toxicity-proxy signals.

3 Model

In this section, we describe the dynamics of the model. We begin by outlining the mechanics of Uniswap v3, followed by a description of how the reference (CEX) market is modeled and how the different cohorts of agents behave. We then introduce three simple dynamic fee schedules: one that adjusts fees in response to CEX or DEX price volatility and one based on a proxy for order-flow toxicity (measured by how much the AMM price is stale during the coupled CEX-DEX dynamics). Finally, we explain how liquidity providers’ PnL is calculated and provide key details of the simulation framework.

3.1 Uniswap v3 Market Mechanism

We consider an AMM that allows swapping between two assets, token 0 and token 1, denoted respectively with x and y . Throughout, prices are quoted as token 1 per unit of token 0. Uniswap v3 implements a concentrated-liquidity AMM [2], meaning that liquidity providers (LPs) choose the price range over which their capital is active. Whenever a swap occurs the liquidity taker (LT) must pay a fee f_t to the LPs. We intentionally put the dependence of the fee on time, since the main objective of this paper is to understand what happens to the economics of the agents if we let the fee have a non trivial dynamics. This section introduces the main objects needed to model Uniswap v3 swaps.

3.1.1 Discrete price grid, ticks, and square-root price

Uniswap v3 discretises the price axis into ticks indexed by integers $i \in \mathbb{Z}$. For simplicity we assume tick spacing $\Delta i = 1$ ¹. The price associated with tick i is

$$p(i) = \gamma^i, \quad \gamma = 1.0001, \quad (1)$$

and Uniswap v3 uses the square-root price because, under constant liquidity, token amounts vary linearly with changes in $s = \sqrt{p}$, which makes swap and position accounting algebraically and computationally simple:

$$s(i) = \sqrt{p(i)} = \gamma^{i/2}. \quad (2)$$

¹Tick spacing is fixed by the pool contract; e.g., for USDC–WETH at 0.05% the tick spacing is 10.

At any time t , the pool is at an active tick i_t and active square-root price s_t . Between two adjacent tick boundaries, the active liquidity is constant and the swap dynamics admit closed-form updates like in Uniswap v2 [1].

3.1.2 Liquidity positions and active liquidity

An LP does not provide liquidity over the entire price axis. Instead, an LP position is specified by a tick interval $[i_a, i_b]$ and a liquidity amount $L^{\text{pos}} > 0$. Define the corresponding boundary square-root prices

$$s_a := s(i_a), \quad s_b := s(i_b). \quad (3)$$

Intuitively, a position is active only when the current price lies inside its chosen range. When active, it contributes depth to the pool and earns a pro-rata share of swap fees; when inactive, it earns no fees and is held entirely in one of the two tokens.

Let L_t denote the active liquidity at time t , i.e., the sum of liquidity across all positions whose ranges contain the active tick:

$$L_t = \sum_{\text{pos}} L^{\text{pos}} \mathbf{1}(i_a^{\text{pos}} \leq i_t < i_b^{\text{pos}}). \quad (4)$$

A useful representation of how liquidity changes across the grid is the liquidity net $\ell(i)$, defined as the net change in active liquidity when crossing tick i :

$$\ell(i) = \sum_{\text{pos}} L^{\text{pos}} \mathbf{1}(i = i_a^{\text{pos}}) - \sum_{\text{pos}} L^{\text{pos}} \mathbf{1}(i = i_b^{\text{pos}}). \quad (5)$$

When the price crosses a tick boundary, some positions become active/inactive and the pool updates L_t by adding $\ell(i)$ at that boundary.

Token composition of a position. Within its active range ($s_a \leq s_t \leq s_b$), the amounts of token 0 and token 1 represented by liquidity L^{pos} at square-root price s_t are

$$x^{\text{pos}}(s_t) = L^{\text{pos}} \left(\frac{1}{s_t} - \frac{1}{s_b} \right), \quad y^{\text{pos}}(s_t) = L^{\text{pos}} (s_t - s_a). \quad (6)$$

Outside the range, the position is entirely in one asset:

$$\begin{aligned} s_t \leq s_a : \quad & x^{\text{pos}} = L^{\text{pos}} \left(\frac{1}{s_a} - \frac{1}{s_b} \right), \quad y^{\text{pos}} = 0, \\ s_t \geq s_b : \quad & x^{\text{pos}} = 0, \quad y^{\text{pos}} = L^{\text{pos}} (s_b - s_a). \end{aligned} \quad (7)$$

These relations highlight why liquidity is “concentrated”: by choosing a narrow interval, an LP increases the effective depth (active liquidity) around that region, at the cost of becoming inactive if the price moves away.

3.1.3 Swap dynamics within a tick

Here we briefly describe how a swap works in Uniswap v3. Consider a swap executed at time t while the price remains within the current tick interval, so that active liquidity L_t is constant. In this regime, Uniswap v3 behaves like a continuous AMM (like Uniswap v2). parameterised by (s_t, L_t) , and the input/output amounts are determined by the change in square-root price.

Let the price move from s to s' with constant liquidity L (we omit the time subscript for readability). The corresponding token amount changes satisfy the standard Uniswap v3 relations:

$$\Delta x = L \left(\frac{1}{s'} - \frac{1}{s} \right), \quad \Delta y = L(s' - s). \quad (8)$$

The sign and direction of the price move depend on the trade direction:

- If a trader swaps token 0 into the pool to receive token 1, the pool's token 0 reserve increases and token 1 reserve decreases, so the price $p = s^2$ moves down (thus $s' < s$). In (8), this corresponds to $\Delta x > 0$ and $\Delta y < 0$.
- If a trader swaps token 1 into the pool to receive token 0, the price moves up (thus $s' > s$), giving $\Delta y > 0$ and $\Delta x < 0$.

Fees are charged on the input side. If the pool fee rate is $f \in (0, 1)$, then for a submitted input amount Δx_{in} only $(1 - f)\Delta x_{\text{in}}$ is effectively used to move the price; the remaining $f\Delta x_{\text{in}}$ is accumulated as fees and later distributed to LPs active over the traversed price region, proportional to their share of liquidity.

3.1.4 Tick crossing and pathwise execution of large swaps

Because liquidity is concentrated, the pool depth varies across prices. A large swap may move the price across multiple ticks, activating and deactivating different LP ranges. Uniswap v3 therefore executes swaps tick by tick: within each tick interval the liquidity is constant and (8) applies, while at tick boundaries the liquidity is updated using the liquidity net.

For definiteness, suppose a trader submits an order of size $\Delta x > 0$ (token 0 in, token 1 out) at time t . Starting from (s_t, i_t, L_t) , the protocol:

1. Computes the maximum amount of token 0 that can be consumed before the next tick boundary is reached (given the current s_t and constant L_t), accounting for fees on the input.
2. If the remaining order is below this capacity, it is fully executed within the current tick by updating s_t using (8), and the swap terminates.
3. Otherwise, the swap moves the price to the tick boundary, consumes the corresponding capacity in that tick, and crosses into the adjacent tick. At the boundary the active tick index and active liquidity are simultaneously updated as

$$i_t \leftarrow i_t + \varepsilon, \quad L_t \leftarrow \begin{cases} L_t + \ell(i_t + 1) & \text{if } \varepsilon = 1 \\ L_t - \ell(i_t) & \text{if } \varepsilon = -1 \end{cases} \quad (9)$$

where ε is the sign of the tick move (-1 for token 0 in and +1 for token 1 in).

4. The procedure repeats until the full order is executed or no active liquidity remains.

This mechanism implies that the swap's price impact is generally piecewise, determined by the sequence of ticks crossed and the corresponding liquidity profile L along the path. Fees are accrued separately within each traversed tick interval and allocated to the LPs whose positions are active over that interval, proportional to their provided liquidity.

3.1.5 Mempool

The simulation evolves in discrete blocks indexed by $t = 1, \dots, T$, mirroring the block-based execution model of Ethereum-like chains. Each block is subdivided into $B \geq 1$ micro-steps (treated as seconds) indexed by $\tau = 1, \dots, B$. Hence, B represents the block time measured in units of a micro-step.

At the end of block t , the simulator records a validated on-chain snapshot

$$S_t = (P_t^{\text{DEX}}, L_t, m_t),$$

where P_t^{DEX} is the AMM price, L_t is the active liquidity state, and m_t is the external CEX mid-price. At the beginning of block $t + 1$, all agents observe S_t and condition their decisions on it. This captures (i) confirmation latency—agents react to the last confirmed on-chain state—and (ii) a common reference for evaluating slippage, basis, and range placement within the block.

During the micro-steps, the external reference price continues to evolve (in our baseline, via a Heston stochastic-volatility specification; see Section 3.2), while the observable on-chain state remains frozen at S_{t-1} from the agents’ perspective. Agents therefore submit actions using information that becomes stale by the time it is executed, consistent with blockchain-style batch settlement.

All intents generated over the B micro-steps are appended to a pending mempool and are not executed immediately. In parallel, an arbitrage intent is constructed to exploit the price discrepancy implied by the frozen on-chain snapshot and the reference at the beginning of the block. At the block boundary, the simulator executes the block via a mempool replay:

1. Collect all intents accumulated during block t .
2. Randomly shuffle the mempool to represent stochastic transaction ordering (i.e. there is no auction for block space and block position).
3. Execute the arbitrage intent with priority (to approximate MEV-style ordering advantages), then execute the remaining intents sequentially against the live AMM state.
4. Record the resulting state as the next validated snapshot S_t .

This mechanism decouples stochastic order arrival from on-chain ordering while keeping per-block targets interpretable and directly controllable through λ_e , and it captures the key latency feature of blockchain systems: agents act on confirmed information, yet execution occurs later under uncertain ordering, with arbitrageurs enjoying systematic priority. [This design abstracts from the full block-building market, but preserves the market-design feature that latency and priority ordering can materially affect volatility, arbitrage, and liquidity provision, as emphasized in both MEV studies and artificial-market analyses of high-frequency trading and cancellation rules \[17, 26, 20, 40\].](#)

3.2 Reference Market (CEX)

The reference market is a centralized exchange (CEX) endowed by an underlying price dynamics and a coupling with the DEX due to the price impact induced by arbitrageurs. We denote by m_t the mid-price at the end of block t . The CEX is used for:

- defining the arbitrage band against the DEX;
- marking LP wealth and fees to market;
- implementing the rebalancing benchmark associated with the Loss-Versus-Rebalancing (LVR).

3.2.1 Heston Volatility Model

To capture stylized facts of asset returns such as volatility clustering and the leverage effect (the negative correlation between returns and volatility changes), we introduce a stochastic volatility regime based on the Heston model [24]. Unlike the geometric Brownian motion (GBM) with constant or deterministically shifting volatility, the variance $v_t = \sigma_t^2$ in this regime follows a mean-reverting stochastic process.

The evolution of the reference market price m_t and its variance v_t is governed by the following system of stochastic differential equations (SDEs):

$$d \ln m_t = \left(\mu - \frac{1}{2} v_t \right) dt + \sqrt{v_t} dW_t^{(1)}, \quad (10)$$

$$dv_t = \kappa(\theta - v_t)dt + \sigma_v \sqrt{v_t} dW_t^{(2)}, \quad (11)$$

where $W_t^{(1)}$ and $W_t^{(2)}$ are standard Brownian motions with correlation ρ , i.e., $\mathbb{E}[dW_t^{(1)} dW_t^{(2)}] = \rho dt$. The parameters are defined as follows:

- μ : The drift of the log-price process.
- θ : The long-run average variance.
- κ : The rate of mean reversion, determining how quickly v_t returns to θ .
- σ_v : The volatility of volatility (vol-of-vol).
- ρ : The correlation parameter. A negative ρ induces the leverage effect, where price drops are associated with increases in volatility.

The long-run variance has been calibrated using data from Binance regarding the ETH/USD pair, between the 1 January 2023 and 31 December 2023.

3.2.2 Market Impact

Trades performed on the CEX are assumed to have zero costs but permanent price impact on the CEX price. Let Δa_τ denote the net signed volume in token 0/token 1 executed on the CEX during microstep τ . Empirically, large trades tend to move prices in a concave way, consistent with the square-root law observed in market-impact studies [37, 29]. Inspired by these findings, the impact function is specified as

$$\text{Impact}_\tau = \eta_{\text{imp}} \cdot \text{sign}(\Delta a_\tau) \sqrt{|\Delta a_\tau|}, \quad (12)$$

with $\eta_{\text{imp}} > 0$ and $\xi \geq 0$.

From the point of view of the arbitrageur, this ensures that large trades move the CEX against the arbitrageur, providing a natural feedback that reduces future mispricing.

3.3 LVR and PnL Accounting

As we already stated, LVR quantifies the structural cost borne by liquidity providers due to the passive nature of Automated Market Makers. It was first introduced by Milionis et al. [32]. This stale-quote cost is the AMM analogue of adverse selection faced by liquidity suppliers in dealer and limit-order-book markets [25, 6, 32, 11]. Unlike active market makers on Centralized Exchanges

(CEX) who can update quotes in real-time, AMM positions offer stale prices that are "picked off" by arbitrageurs whenever the external market moves. LVR measures the cumulative value extracted by these arbitrageurs—or equivalently, the opportunity cost of providing liquidity compared to a specific benchmark strategy.

To isolate the cost of liquidity provision from market risk, the authors in [32] compare the LP's wealth to a theoretical rebalancing strategy. This benchmark strategy is defined by two key properties:

1. **Inventory Matching:** At every instant t , the rebalancing strategy holds the exact same quantity of the risky asset (token 0), denoted by x_t , as the actual LP position in the AMM.
2. **Frictionless Execution:** Whenever the inventory x_t changes (due to trading activity in the pool), the rebalancing strategy executes the corresponding trade at the fresh, external reference market price m_t , rather than at the stale AMM price P_t^{DEX} .

Because the rebalancing strategy executes the same sequence of trades as the AMM but at better prices (selling at $m_t > P_t^{DEX}$ when prices rise, and buying at $m_t < P_t^{DEX}$ when prices fall), its value V_t^{reb} strictly exceeds the LP's value V_t^{LP} . The difference is the LVR:

$$LVR_t = V_t^{reb} - V_t^{LP}. \quad (13)$$

Milionis et al. [32] show that if the external reference price m_t follows a diffusion process with instantaneous volatility σ_t , LVR accumulates as a function of the price volatility and the AMM's marginal liquidity. The cumulative LVR over a period $[0, T]$ is given by:

$$LVR_{[0,T]} = \int_0^T \frac{\sigma_t^2 m_t^2}{2} |x^{*'}(m_t)| dt, \quad (14)$$

where $x^*(p)$ is the AMM's demand curve (the inventory of token 0 held by the pool at price p), and $|x^{*'}(p)| = |\frac{dx}{dp}|$ is the marginal liquidity, representing how aggressively the AMM trades against price changes.

In the specific context of Uniswap v3, the inventory function is determined by the active concentrated liquidity L_t . Within a specific tick range, the relationship between token 0 inventory and square-root price $s = \sqrt{p}$ is given by Equation (7):

$$x(s) = L_t \left(\frac{1}{s} - \frac{1}{s_b} \right). \quad (15)$$

To apply the general LVR Equation 14, we compute the marginal liquidity with respect to the price $p = s^2$. Using the chain rule:

$$\left| \frac{dx}{dp} \right| = \left| \frac{dx}{ds} \frac{ds}{dp} \right| = \left| \left(-\frac{L_t}{s^2} \right) \left(\frac{1}{2\sqrt{p}} \right) \right| = \frac{L_t}{2p\sqrt{p}}. \quad (16)$$

Substituting this marginal liquidity back into Equation 14 (with $p = m_t$), we obtain the instantaneous LVR specific to a Uniswap v3 position:

$$l_t^{v3} = \frac{\sigma_t^2 m_t^2}{2} \cdot \frac{L_t}{2m_t\sqrt{m_t}} = \frac{\sigma_t^2 L_t \sqrt{m_t}}{4}. \quad (17)$$

Equation 17 provides the key intuition for our agent-based model: the cost of providing liquidity scales linearly with the active liquidity L_t and the price level $\sqrt{m_t}$, and quadratically with volatility σ_t . Crucially, because L_t is the concentrated liquidity, LVR is only incurred when the reference price m_t falls within the LP's active tick range (where $L_t > 0$); inactive positions do not suffer LVR.

3.3.1 Rebalancing Benchmark in Discrete Time

The ABM computes LVR using a discrete-time rebalancing benchmark that faithfully tracks the LP's token exposure.

Fix an LP and let m_t be the CEX mid at the end of block t . Let x_t denote the LP's total exposure in token 0 at time t (the sum over all its AMM positions) and y_t its cash holdings in token 1. The LP's mark-to-market wealth is

$$V_t^{\text{LP}} = x_t m_t + y_t. \quad (18)$$

We construct a benchmark portfolio $(x_t^{\text{reb}}, y_t^{\text{reb}})$ with the following properties:

- it always holds the same exposure in token 0 as the LP,

$$x_t^{\text{reb}} = x_t; \quad (19)$$

- it trades only at the CEX mid m_t ;
- it is self-financing.

Let the benchmark be initialized at $t = 0$ with

$$x^{\text{reb}} = x, \quad y_0^{\text{reb}} = y_0, \quad V_0^{\text{reb}} = V_0^{\text{LP}}. \quad (20)$$

Then we update it in two ways:

Price moves (passive holding). When the CEX price moves from m_{t-1} to m_t while the benchmark holds x_{t-1}^{reb} units of token 0, the benchmark's value changes by

$$\Delta V_t^{\text{reb,price}} = x_{t-1}^{\text{reb}}(m_t - m_{t-1}), \quad (21)$$

so that

$$V_t^{\text{reb}} = V_{t-1}^{\text{reb}} + x_{t-1}^{\text{reb}}(m_t - m_{t-1}). \quad (22)$$

Rebalancing trades (active adjustments). Whenever the LP's exposure changes from x_{t-}^{LP} to x_{t+}^{LP} due to AMM interactions (minting, burning, price moves within the pool), the benchmark trades at the CEX price m_t to match the new exposure:

$$x_{t+}^{\text{reb}} = x_{t+}^{\text{LP}}, \quad y_{t+}^{\text{reb}} = y_{t-}^{\text{reb}} - (x_{t+}^{\text{reb}} - x_{t-}^{\text{reb}})m_t. \quad (23)$$

This ensures the benchmark is self-financing: changes in x^{reb} are funded by opposite changes in y^{reb} , with no external cash injections.

By construction, the benchmark value is

$$V_t^{\text{reb}} = x_t^{\text{reb}} m_t + y_t^{\text{reb}}. \quad (24)$$

3.4 Agents

The ecosystem comprises several types of agents with distinct objectives, information sets, and constraints. In this section, we describe their nature and behaviour.

3.4.1 Arbitrageurs

Arbitrageurs try to keep the price discrepancy between the DEX and the CEX price as small as possible, up to wedges induced by the AMM taker fee and flash-loan funding. At the beginning of each block they observe the validated snapshot (m_t, P_t^{DEX}) , where m_t is the CEX mid-price (token 1 per token 0) and P_t^{DEX} is the current marginal DEX price.

Let $f_t \in [0, 1)$ denote the taker fee at time t and $\phi_{\text{flash}} \geq 0$ the flash-loan fee, interpreted as a proportional cost on the borrowed principal. Arbitrage is modeled as: (i) borrow the input token via flash loan; (ii) swap against the AMM; (iii) unwind on the CEX at price m_t to repay principal plus fee. All payoffs are measured in token 1 units. Because the AMM price moves during execution, profitability is inherently size-dependent. The arbitrageur therefore chooses an executed swap size $q \geq 0$ to maximize net profit. [The arbitrageur is therefore both the mechanism enforcing price alignment and the channel through which stale AMM quotes become realized LP losses \[5, 32, 19, 10\].](#) We have two possible situations, described below.

DEX overpriced: token 0 in. If P_t^{DEX} is high relative to m_t , the arbitrageur borrows token 0, inputs q token 0 into the AMM, receives an output $Y_t(q)$ in token 1, and repays $(1 + \phi_{\text{flash}})q$ token 0 by buying it on the CEX at price m_t . Net profit is

$$\Pi_t^{(0 \rightarrow 1)}(q) = Y_t(q) - m_t(1 + \phi_{\text{flash}})q - c_{\text{fix}}, \quad (25)$$

where $c_{\text{fix}} \geq 0$ captures fixed costs (e.g. gas) in token 1 units. For simplicity, we will consider $c_{\text{fix}} = 0$.

Fees are charged on the input: only the fee-adjusted amount $u = (1 - f_t)q$ moves the AMM state. Let $P_t(q)$ denote the marginal DEX price (token 1 per token 0) after executing size q (piecewise across ticks). For an infinitesimal increase dq , the AMM output satisfies

$$\frac{dY_t(q)}{dq} = (1 - f_t)P_t(q). \quad (26)$$

Hence the marginal net profit is

$$\frac{d\Pi_t^{(0 \rightarrow 1)}(q)}{dq} = (1 - f_t)P_t(q) - m_t(1 + \phi_{\text{flash}}), \quad (27)$$

and an interior optimum $q^* > 0$ satisfies the first-order condition

$$P_t(q^*) = \frac{m_t(1 + \phi_{\text{flash}})}{1 - f_t}. \quad (28)$$

Operationally, this means the arbitrageur trades token 0 into the AMM until the marginal DEX price is pushed down to the boundary (28), or until liquidity is exhausted.

DEX underpriced: token 1 in. If P_t^{DEX} is low relative to m_t , the arbitrageur borrows token 1, inputs q token 1 into the AMM, receives an output $X_t(q)$ in token 0, sells it on the CEX for $m_t X_t(q)$ token 1, and repays $(1 + \phi_{\text{flash}})q$ token 1. Net profit is

$$\Pi_t^{(1 \rightarrow 0)}(q) = m_t X_t(q) - (1 + \phi_{\text{flash}})q - c_{\text{fix}}. \quad (29)$$

Within the AMM, the fee-adjusted input is again $u = (1 - f_t)q$. Let $P_t(q)$ denote the marginal DEX price after executing size q (now increasing in q). The marginal token 0 output per unit token 1 input is the reciprocal price, so

$$\frac{dX_t(q)}{dq} = (1 - f_t) \frac{1}{P_t(q)}. \quad (30)$$

Thus

$$\frac{d\Pi_t^{(1 \rightarrow 0)}(q)}{dq} = m_t(1 - f_t) \frac{1}{P_t(q)} - (1 + \phi_{\text{flash}}), \quad (31)$$

and an interior optimum $q^* > 0$ satisfies

$$P_t(q^*) = \frac{m_t(1 - f_t)}{1 + \phi_{\text{flash}}}. \quad (32)$$

That is, the arbitrageur trades token 1 into the AMM until the marginal DEX price is pushed up to the boundary (32), or until liquidity is exhausted.

No-arbitrage band and execution rule Equations (32)–(28) imply a two-sided no-arbitrage band (in marginal-price terms):

$$\frac{m_t(1 - f_t)}{1 + \phi_{\text{flash}}} \leq P_t^{\text{DEX}} \leq \frac{m_t(1 + \phi_{\text{flash}})}{1 - f_t}. \quad (33)$$

If P_t^{DEX} lies inside (33), the arbitrageur does not trade. Otherwise, it computes the executed size q^* by simulating the AMM swap (accounting for the input fee) until the marginal DEX price reaches the nearest boundary (32) or (28), or until no further progress is possible due to finite liquidity. The trade is executed only if the corresponding net profit $\Pi_t^{(0 \rightarrow 1)}(q^*)$ or $\Pi_t^{(1 \rightarrow 0)}(q^*)$ is strictly positive, in which case flash-loan costs (and fixed costs) generate a non-zero region of persistent mispricings. The band in 33 is similar to the one derived in [5] and it becomes exactly equal if $\phi_{\text{flash}} = 0$. Our setting, however, improves the realism of these band by introducing frictions via the flash-loan fees and finite liquidity.

Execution and Impact The arbitrageur trades against the pool to move the DEX price as close as possible to the relevant band boundary in (33), subject to liquidity depth. It may not always be possible to reach the band exactly; in that case, the arbitrageur consumes all available depth in the profitable direction and stops.

Arbitrage trades are assumed to incur CEX impact as in (12). In the baseline regime, apart from this permanent impact and the flash-loan fee, arbitrageurs are:

- well-capitalized (no additional capital constraint),
- continuously monitoring (one attempt every block),
- executed before other mempool flow in block mode.

This creates a harsh but realistic environment for LPs, as any significant mispricing that remains profitable after funding costs is rapidly harvested.

3.4.2 Liquidity Providers

LPs deposit liquidity into the AMM and collect a pro-rata share of swap fees whenever their positions are active. We model three LP cohorts: (i) *passive* LPs, who provide wide ranges and adjust infrequently; (ii) *active* LPs, who concentrate liquidity near the current price and manage positions more aggressively; and (iii) a MEV searcher performing *Just-In-Time* (JIT) liquidity provision (the *Jiter*).

LP state Each (non-JIT) LP j maintains:

- a cash wallet $\mathcal{W}_{j,t}^1$ held entirely in token 1;
- a set of open positions composed by the lower tick i_a , upper tick i_b , liquidity deposited L^{pos} in $[i_a, i_b] \implies \mathcal{P}_{j,t} = \{(i_a, i_b, L^{\text{pos}})_t\}$;
- internal review and cooldown clocks governing when the LP can revise positions; the review clock determines when the LP next reconsiders its range, while the cooldown clock enforces a waiting period after a burn before a new mint can be submitted;
- an associated rebalancing benchmark used for LVR accounting.

LPs are *cash-budgeted*: between actions they do not carry token 0 in their wallet. When minting a new position, the LP uses token 1 cash to acquire (conceptually, on the CEX) the exact token 0 amount required by Uniswap v3 for the chosen range at the prevailing CEX price m_t , and deposits the resulting pair of token amounts into the AMM. When burning, the LP receives the position's underlying token amounts and accrued fees, and then immediately converts any token 0 proceeds back into token 1 at price m_t . These conversions occur at the CEX price without explicit transaction fees but they induce permanent price impact.

Minting a position Consider a mint decision at block t for LP j with tick range $[i_a, i_b]$ (with $i_a < i_b$) and liquidity $L > 0$. Let $s_t = \sqrt{P_t^{\text{DEX}}}$, $s_a = s(i_a)$, $s_b = s(i_b)$, where $s(\cdot)$ maps ticks to square-root prices (Section 3.1). By the Uniswap v3 algebra, minting liquidity L over $[s_a, s_b]$ requires depositing $(\Delta x(L), \Delta y(L))$, where

$$(\Delta x(L), \Delta y(L)) = \begin{cases} \left(L \left(\frac{1}{s_a} - \frac{1}{s_b} \right), 0 \right), & s_t \leq s_a, \\ \left(L \left(\frac{1}{s_t} - \frac{1}{s_b} \right), L(s_t - s_a) \right), & s_a < s_t < s_b, \\ (0, L(s_b - s_a)), & s_t \geq s_b. \end{cases} \quad (34)$$

A key property is that, once $(s_t; s_a, s_b)$ is fixed, both $\Delta x(L)$ and $\Delta y(L)$ scale linearly with L . Therefore, valuing token 0 at the CEX price m_t , the LP can mint at maximum a liquidity $L_{j,t}^{\max}([i_a, i_b])$ such that

$$\mathcal{W}_{j,t}^1 = L_{j,t}^{\max}([i_a, i_b]) \left(m_t \Delta x(1) + \Delta y(1) \right), \quad (35)$$

where $(\Delta x(1), \Delta y(1))$ denotes the deposit required for one unit of liquidity in the same state $(s_t; s_a, s_b)$. Hence

$$L_{j,t}^{\max}([i_a, i_b]) = \frac{\mathcal{W}_{j,t}^1}{m_t \Delta x(1) + \Delta y(1)}. \quad (36)$$

To model heterogeneous aggressiveness while avoiding all the mints to be “all-in”, each time the LP draws a wallet utilization factor $\eta_{j,t} \in (0, 1]$:

$$Z_{j,t} \sim \log \mathcal{N}(\mu_{\text{mint}}, \sigma_{\text{mint}}), \quad \eta_{j,t} = \min\{1, Z_{j,t}\}. \quad (37)$$

The executed mint liquidity is

$$\Delta L_{j,t}^{\text{exec}} = \eta_{j,t} L_{j,t}^{\max}([i_a, i_b]). \quad (38)$$

If $\Delta L_{j,t}^{\text{exec}} = 0$, the mint is skipped. Otherwise, the LP uses token 1 cash to finance the mint: it acquires on the CEX, at price m_t , the required token 0 amount $\Delta x(\Delta L_{j,t}^{\text{exec}})$, contributes the token 1

amount $\Delta y(\Delta L_{j,t}^{\text{exec}})$ from its cash wallet, and deposits $(\Delta x(\Delta L_{j,t}^{\text{exec}}), \Delta y(\Delta L_{j,t}^{\text{exec}}))$ into the AMM. The cash wallet is debited by the token 1 value of this deposit:

$$\mathcal{W}_{j,t+}^1 = \mathcal{W}_{j,t}^1 - \left(m_t \Delta x(\Delta L_{j,t}^{\text{exec}}) + \Delta y(\Delta L_{j,t}^{\text{exec}}) \right), \quad (39)$$

and the position $(i_a, i_b, \Delta L_{j,t}^{\text{exec}})$ is added to $\mathcal{P}_{j,t}$. Any unspent budget remains in token 1; in particular, the LP does not retain a residual token 0 inventory.

Burning a position When LP j burns a position $p = (i_a, i_b, L^{\text{pos}}) \in \mathcal{P}_{j,t}$ at state s_t , the LP receives (i) the position's current underlying token amounts $(\Delta y(\Delta L_{j,t}^{\text{exec}}), \Delta x(\Delta L_{j,t}^{\text{exec}}))$ and (ii) its accrued fees $(F_{1,t}^p, F_{0,t}^p)$. The underlying amounts are given by the same v3 formulas as (34) with $L = L^{\text{pos}}$. Under cash budgeting, token 0 proceeds are immediately converted into token 1 at price m_t , so the wallet updates as

$$\mathcal{W}_{j,t+}^1 = \mathcal{W}_{j,t}^1 + \left(\Delta y(\Delta L_{j,t}^{\text{exec}}) + F_{1,t}^p \right) + m_t \left(\Delta x(\Delta L_{j,t}^{\text{exec}}) + F_{0,t}^p \right). \quad (40)$$

Fees are accrued tick-by-tick during swap execution: in each traversed tick interval, the input-side fee flow is allocated to the LPs active in that interval proportionally to their provided liquidity L^{pos} .

Passive LPs Passive LPs provide baseline liquidity for swaps. Their strategy is:

- **Wide ranges:** Passive LPs mint positions symmetrically around the current DEX price, setting the range width as a fixed percentage of the entry price.
- **Exogenous activity:** Passive mint and burn decisions are probabilistic and do not depend on microstructure signals.
- **Token-budgeted sizing:** Conditional on minting a range, the position size is determined by (35)–(38), i.e. by available token inventories and the utilization factor.

Active LPs Active LPs provide more liquidity near the current price and manage positions dynamically. [This active cohort is motivated by empirical evidence that Uniswap v3 LP outcomes depend strongly on range width and position timing, and by strategic-liquidity models in which LPs dynamically adjust ranges to optimize fee income and inventory exposure \[23, 18\].](#) They differ from passive LPs along several dimensions:

- **Adaptive width:** The width w_t (in ticks) of each active position is determined from an EWMA of a signal related to recent volatility, plus a binomial noise term to introduce heterogeneity:

$$w_t = \min \left(\max(\tilde{w}_t + \eta_t, w_{\min}), w_{\max} \right), \quad (41)$$

where $\tilde{w}_t$ is a deterministic function of the chosen signal, η_t is a random width perturbation, and $[w_{\min}, w_{\max}]$ are bounds.

- **Re-centering:** Each active position tracks the number of consecutive blocks for which the DEX price has remained outside its range. If this counter exceeds a threshold $k_{\text{out}} \sim U[k_{\text{out}}^{\min}, k_{\text{out}}^{\max}]$, the LP enqueues a recenter intent: burn the old position and mint a new one centered at the current reference, with width w_t . The new mint is executed subject to token feasibility (35)–(38).

- **Take-profit / stop-loss:** Active LPs monitor the PnL of each position relative to a HODL benchmark of the initial tokens, meaning the value the LP would obtain by simply holding the two tokens that were initially required to mint the position and marking them to the external reference price. If the position’s PnL exceeds a take-profit threshold θ_{TP} or falls below a stop-loss threshold $-\theta_{\text{SL}}$, the LP exits by burning the position.

Active LPs therefore tend to trade more frequently and run narrower ranges, making them more exposed to LVR but potentially able to harvest higher fee density.

Review clocks and cooldowns Active and passive LPs are endowed with a review clock that governs how often they can revise positions. Inter-review times are drawn from a geometric distribution with mean τ_{review} , so that at any block roughly a fraction $1/\tau_{\text{review}}$ of LPs are “due” to reconsider their ranges. When an LP burns a position, it enters a cooldown during which it cannot mint new positions. This mechanism introduces operational latency, mitigates high-frequency strategic oscillations and prevents synchronization among all the LPs.

JIT LP (Jiter) The Jiter is a MEV-style liquidity provider, motivated by documented JIT liquidity and priority-ordering behavior in AMMs [38, 17], that tries to capture swap fees by providing liquidity only for the short window around a large swap. In each block, it participates with probability p_{JIT} and, when it does, it targets the single largest pending swap by input notional (valued in token1). When the target reaches execution, the swap is wrapped as $\text{jit mint} \rightarrow \text{swap} \rightarrow \text{jit burn}$: the Jiter mints a one-tick-wide position, lets the swap trade against this added liquidity, and then immediately burns the position. The mint size is not a fixed quantity; instead, it is chosen by trading off the expected fee income the Jiter can earn inside that tick (which depends on how much of the swap remains in-range and on the Jiter’s share of tick liquidity) against the cost of flash-financing the minted principal. The size is capped so that the Jiter cannot exceed a prescribed fraction of the tick’s active liquidity after minting (90%), and the mint is skipped whenever the resulting expected fee capture does not cover the flash-loan fee. We model the mint as flash-funded and charge a per-mint financing cost given by the same fee the arbitrage faces for his flash loan (0.01%); upon burning, it withdraws principal plus accrued fees and nets any remaining token0 exposure by converting it back to token1 at the current reference price.

PnL Decomposition Let F_t denote the cumulative fees earned by the LP up to time t , valued in token 1 using the CEX price m_t .

The key identity enforced by the model is

$$V_t^{\text{LP}} = V_t^{\text{reb}} + F_t - \text{LVR}_t, \quad (42)$$

where LVR_t is the cumulative loss-versus-rebalancing experienced by the LP up to time t . Rearranging,

$$\text{LVR}_t = V_t^{\text{reb}} - V_t^{\text{LP}} + F_t. \quad (43)$$

From (42) we obtain the following PnL notions:

- **Unhedged PnL** (full economic outcome):

$$\Pi_t^{\text{unhedged}} = V_t^{\text{LP}} - V_0^{\text{LP}}. \quad (44)$$

This includes all market beta and directional exposure.

- Hedged PnL (liquidity-provision economics only):

$$\Pi_t^{\text{hedged}} = V_t^{\text{LP}} - V_t^{\text{reb}} = F_t - \text{LVR}_t. \quad (45)$$

This measures how much of the fees survive after subtracting adverse selection losses. In the case of the Jiter, Π_t^{hedged} is reduced also by the flash loan cost.

By design, hedged PnL cannot be improved by taking more price risk: any strategy that alters the LP's exposure path x_t changes both the benchmark and the AMM wealth in lockstep. Only the balance between fees collected and LVR determines Π_t^{hedged} .

3.4.3 Smart router

To endogenize venue choice, we introduce a smart-routing trader that allocates each potential trade between the automated market maker (AMM, i.e., the DEX) and the reference centralized exchange (CEX) according to a best-execution criterion. The inclusion of a smart router is central to the model, as it captures the endogenous relative attractiveness of AMM versus CEX execution. When the liquidity-taker fee on the AMM is high, the smart router optimally favors execution on the CEX, reducing the volume routed to the AMM and, consequently, the fee income accruing to liquidity providers (LPs). Conversely, when the liquidity-taker fee is low, the smart router preferentially routes trades to the AMM, increasing fee revenues for LPs. This mechanism highlights how the ratio of CEX-routed to AMM-routed trades emerges as an informative metric of DEX efficiency

Smart router draws a token 1 notional $Y^{\text{not}} \sim \exp(\mathcal{N}(\mu_{\text{tr}}, \sigma_{\text{tr}}^2))$ and a direction side $\in \{X \rightarrow Y, Y \rightarrow X\}$ with equal probability. For $X \rightarrow Y$ (sell token 0), the smart router fixes a token 0 input $\Delta x^{\text{int}} = Y^{\text{not}}/m_t$ and queries the AMM quote Δy^{DEX} . The CEX benchmark output is $\Delta y^{\text{CEX}} = \Delta x^{\text{int}} m_t$ corresponding to executing the same sale at price m_t . The best-execution requirement is

$$\Delta y^{\text{DEX}} \geq \Delta y^{\text{CEX}}, \quad (46)$$

If (46) fails, the trade is routed entirely to the CEX; otherwise, the AMM leg is submitted to the mempool. For $Y \rightarrow X$ (buy token 0), the trader fixes a token 1 input $\Delta y^{\text{int}} = Y^{\text{not}}$, queries Δx^{DEX} and compares it to the CEX benchmark $\Delta x^{\text{CEX}} = \Delta y^{\text{int}}/m_t$ via

$$\Delta x^{\text{DEX}} \geq \Delta x^{\text{CEX}}. \quad (47)$$

At execution time (when the mempool is replayed), the engine re-quotes the AMM against a baseline computed from the last validated DEX reference price P_t^{ref} . The AMM execution is skipped if the realized quote violates a relative slippage tolerance, i.e. if $\Delta y^{\text{act}}/\Delta y^{\text{base}} < 1 - \delta_{\text{slip}}$ (or analogously $\Delta x^{\text{act}}/\Delta x^{\text{base}} < 1 - \delta_{\text{slip}}$), where $\delta_{\text{slip}} \in (0, 1)$ is configured. This agent therefore provides a parsimonious micro-founded mechanism for fee-sensitive flow: as the AMM fee increases or liquidity thins, a larger fraction of smart-router volume migrates to the CEX². His PnL is evaluated according to the CEX price, hence his PnL will depend mainly on the discrepancies between our AMM and the reference market. The more the DEX is attractive, the higher the smart router's PnL will be. Finally, CEX trades experience permanent impact as explained in Section 3.2.2.

²It is worth noting, however, that in practice a trader may choose to route an order to an AMM not solely because it offers the best immediate price, but as part of a broader strategy involving other liquidity pools or DeFi protocols. Such behavior can be captured by relaxing the best-execution assumption and introducing a small tolerance around the optimal price.

3.4.4 Noise Trader

The noise trader represents uninformed or urgency-driven flow:

- Her trade sizes are drawn from a log-normal distribution in notional units of token 1, with parameters $(\mu_{\text{trade}}, \sigma_{\text{trade}})$.
- Trade direction (buy vs. sell) is chosen at random with equal probability.
- Noise trader ignores the CEX price m_t and thus have no valuation discipline.
- She enforces a slippage constraint: if the implied execution price on the DEX deviates from their reference quote by more than a tolerance (relative to the reference), the trade is aborted ³.

Noise traders pay fees to LPs and add volume, but their trades also move the AMM price away from the CEX mid, creating opportunities for arbitrageurs and, hence, contributing to LVR. The noise trader’s PnL is computed at the end of each block using the CEX price to convert everything in token 1.

Micro-step event probabilities. As already said in Section 3.1.5, the simulator advances in discrete blocks (indexed by $t = 1, \dots, T$). Each block contains B micro-steps (indexed by $\tau = 1, \dots, B$) that represent within-block order arrival, while execution is aggregated at the block boundary via a mempool replay. For each event class e (e.g., noise trades, LP mints, LP burns), the configuration specifies a micro-step arrival probability $p_e \in [0, 1]$.

Within each block, arrivals are sampled independently across micro-steps as

$$A_{e,t,\tau} \sim \text{Bernoulli}(p_e), \quad \tau = 1, \dots, B,$$

and one event of type e is instantiated whenever $A_{e,t,\tau} = 1$. The resulting number of arrivals within block t is

$$N_{e,t} = \sum_{\tau=1}^B A_{e,t,\tau} \sim \text{Binomial}(B, p_e),$$

hence $\mathbb{E}[N_{e,t}] = Bp_e$ and the expected number of events per block scales linearly with the block time B .

Each realized event is instantiated as an intent and assigned to an eligible agent (or position) drawn from the corresponding candidate set (e.g., agents that are “due” and not in cooldown), allowing multiple intents per agent in a block when applicable. All intents accumulated over the block are then executed at the block boundary (mempool replay), which decouples stochastic arrival from on-chain ordering while keeping the arrival process directly controllable through p_e .

3.5 Dynamic Fee Schedules

A central feature of the model is a time-varying taker fee $f_t \in [0, 1)$, representing the fraction of notional charged to traders and paid to LPs. At each step, a raw fee proposal f_t^{raw} is computed from a chosen signal and then passed through a common controller that enforces bounds, hysteresis, rate limits, and cooldown. [In market-design terms, this controller treats the AMM fee as an endogenous transaction-cost rule, analogous to fee and tax interventions studied in artificial financial markets, but here targeted at AMM adverse selection \[39, 30, 35, 11\].](#)

We consider three modes, based on:

³This behaviour mimics the default slippage tolerance set using the Uniswap v3 web interface.

- i) static
- ii) volatility
- iii) toxicity

EWMA smoothing Several controllers smooth a noisy observation process x_t through an exponentially weighted moving average (EWMA)

$$v_t = \lambda v_{t-1} + (1 - \lambda) x_t, \quad (48)$$

with $\lambda \in (0, 1)$. In the implementation, λ is parameterized by a half-life in steps h ,

$$\lambda = \exp\left(-\frac{\ln 2}{h}\right), \quad \Rightarrow \quad 1 - \lambda \text{ is the EWMA gain.} \quad (49)$$

Static fee The fee is constant:

$$f_t = f_0, \quad \forall t, \quad (50)$$

where f_0 is the initial fee level (and the fixed fee in static mode).

Volatility based fee This controller reacts to realized volatility of the price. We employ two variants according to the venue we take as reference. For simplicity, we report here the specification relative to the CEX price. It is totally similar in the case of the (validated) DEX price⁴. Define

$$vol_obs_t = (\ln(m_t) - \ln(m_{t-1}))^2, \quad (51)$$

and its EWMA $\hat{\sigma}_t = \text{EWMA}(vol_obs_t)$ as in (48). The raw fee proposal is:

$$f_t^{\text{raw}} = k_\sigma \sqrt{\hat{\sigma}_t}, \quad (52)$$

where $k_\sigma > 0$ scales sensitivity.

Toxicity based fee This controller measures arbitrage “toxicity” via the price gap that exceeds the current fee band. Let f_{current} be the fee currently applied by the pool and define the fee band in log-space:

$$f_{\text{bandln}} = -\ln(1 - f_{\text{current}}). \quad (53)$$

Let the absolute log price gap be

$$\Delta_{\text{cex_dex}} = |\ln(P_t^{\text{DEX}}) - \ln(m_t)|. \quad (54)$$

The excess basis observation (using the *current* fee) is

$$B_{\text{obs},t} = \max(0, \Delta_{\text{cex_dex}} - f_{\text{bandln}}), \quad (55)$$

which is then smoothed as $B_{\hat{t}} = \text{EWMA}(B_{\text{obs},t})$. Converting to ticks using the log tick size $\Delta i \log \gamma$,

$$\tilde{B}_t = \frac{B_{\hat{t}}}{\Delta i \log \gamma}. \quad (56)$$

The raw fee is:

$$f_t^{\text{raw}} = k_{\text{tox}} \tilde{B}_t, \quad (57)$$

with $k_{\text{tox}} > 0$.

⁴Using the CEX price requires the AMM to have an oracle able to read that price. In practice, this is possible but it adds a new risk relative to oracle failures.

Shared fee update controller Regardless of mode, f_t^{raw} is transformed into the applied fee through the same mechanism:

(i) Define the increment

$$\Delta f_t = f_t^{\text{raw}} - f_t \quad (58)$$

(ii) Hysteresis and step-size limits

If the change is too small, no update is performed. If it is too high, it is capped. Let $\Delta f_{\max}, \Delta f_{\min}$ be the max and min possible fee updates:

$$|\Delta f_t| < \Delta f_{\min} \Rightarrow \text{no update.} \quad (59)$$

$$|\Delta f_t| > \Delta f_{\max} \Rightarrow \Delta f_t = \Delta f_{\max} \quad (60)$$

(iii) Fee update

The new fee is set to

$$f_{t+1} = f_t + \Delta f_t \quad (61)$$

clamped between a minimum and maximum fee value to prevent extreme situations.

(iv) Cooldown

After a fee change is scheduled, further changes are ignored for τ_{cool} steps, preventing rapid oscillations. Overall, this controller keeps fees bounded, smooth, and responsive only to meaningful shifts in the underlying signal.

This controller ensures that the fee path is bounded, piecewise smooth, and only reacts to material changes in the underlying signals.

3.6 Initialization

The model is initialized by setting an initial pool price P_0^{DEX} and seeding the tick grid with a baseline liquidity profile, distributed as a binomial mass concentrated around the initial DEX price. This “liquidity hill” is allowed to evolve endogenously as swaps move the active price across ticks, but the positions that constitute it are assigned to a separate LP cohort that is excluded from the PnL accounting. Its sole purpose is to provide initial depth so that early trades can execute without pathological price jumps, thereby ensuring a well-defined transient phase before the endogenous LP population and trading flows drive the system toward its stationary regime.

4 Results

In this section we illustrate the results obtained with the ABM framework introduced in the previous sections. The goal is to assess whether dynamic fees can mitigate adverse selection (LVR) under block-time execution, and to quantify the trade-offs with competitiveness (DEX share) and with strategic liquidity provision (active LPs and JIT liquidity).

We consider three models with increased complexity:

1. **Model 0:** arbitrageur, noise trader, smart router, and a passive LP cohort only. This baseline isolates block-time stale-quote risk without strategic liquidity provision.
2. **Model 1:** we introduce an active (more concentrated) LP cohort alongside passive LPs (equal split of the LP population) and compare hedged PnL across cohorts under different fee policies.
3. **Model 2:** we additionally allow a JIT liquidity provider (the “Jiter”) to operate, to study how MEV-style liquidity affects the incidence of fees and LVR across LP cohorts.

All simulations share a common set of parameters (Table 1). As stated in the introduction, our objective is not to construct a data-generating process that accurately reproduces empirical market properties, but rather to develop a framework that allows us to study how different types of agents interact and how the introduction of dynamic fee mechanisms affects the economics of liquidity provision. Accordingly, parameter values are chosen to generate plausible and interpretable dynamics, providing a meaningful setting in which qualitative reasoning and comparative analysis can be carried out. Several parameters (e.g., the order-flow intensity per block) are chosen to be consistent with the empirical calibration in [34]. We anyway provides a brief section in which we shows that even without a proper empirical calibration, the model is able to generate realistic price and trade dynamics, see Section 4.1.

Unless otherwise stated, summary figures report mean $\pm$ standard error across 100 independent seeds per (model, fee schedule) configuration.

The main metrics of interest are:

- the **hedged LP PnL** (Section 3.4.2), which quantifies the profitability of liquidity provision in the AMM net of market risk:

$$\Pi_T^{\text{hedged}} = F_T - \text{LVR}_T - \mathbb{K}_{JIT} C_{\text{flash}}, \quad (62)$$

where $\mathbb{K}_{JIT} = 1$ for the Jiter cohort and 0 otherwise;

- the **fraction of trades routed to the DEX** relative to total trades (DEX share), which serves as a measure of the AMM’s attractiveness compared to the CEX:

$$\text{DEX}_{\text{share}} = \frac{\# \text{DEX trades}}{\# \text{CEX+DEX trades}} \quad (63)$$

- the **liquidity-taker fee level** f_t . Holding all other parameters fixed, higher fees increase LP revenues but reduce the competitiveness of the AMM, thereby discouraging order routing by the smart router in favor of the CEX.

Symbol	Description	Value
Timeline & horizon		
τ	Block duration (mempool length), in seconds / micro-steps per block	5
T	Total number of simulation blocks	13000
T_{skip}	Blocks skipped to reach equilibrium	3000
Reference market (CEX): Heston		
μ	Drift of $\ln m_t$ per micro-step	0
σ_0	Initial instantaneous volatility per micro-step ($v_0 = \sigma_0^2$)	1.5×10^{-4}
κ	Mean reversion speed of variance v_t	1.0
θ	Long-run variance level	1.0×10^{-8}
σ_v	Volatility of variance (vol-of-vol)	1.0×10^{-3}
ρ	Corr. between price and variance shocks	-0.5
v_0	Initial variance level	2.25×10^{-8}
Noise traders		
μ_{trade}	Log-normal location for trade notional (token 1 units)	2.5
σ_{trade}	Log-normal scale for trade notional (token 1 units)	1.5
δ_{slip}	Max tolerated relative slippage (trade aborted if exceeded)	0.01
Arbitrageurs		
ϕ_{flash}	Flash-loan funding cost fraction	1.0×10^{-4}
Liquidity providers (LPs)		
N_{LP}	Number of LP agents	10
$w_{\%}^{\text{pass}}$	Passive LP total width ($\pm w_{\%}^{\text{pass}}/2$ around reference price)	5%
τ_{review}	Mean LP review cadence (blocks)	5
w_{min}	Minimum programmatic mint width (ticks)	10
w_{max}	Maximum programmatic mint width (ticks)	1.77454×10^6
h_{basis}	Half-life (blocks) of EWMA basis signal driving LP widths	1
s	Width sensitivity (ticks per unit basis signal)	1
$n_{\text{bin}}, p_{\text{bin}}$	Binomial noise on widths: trials and success probability	(70, 0.5)
$k_{\text{out}}^{\text{min}}, k_{\text{out}}^{\text{max}}$	Out-of-range blocks before recenter threshold, $k_{\text{out}} \sim U[\cdot]$	(10, 20)
$\mu_{\text{mint}}, \sigma_{\text{mint}}$	Log-normal parameters for mint-size scaling	(-1, 1.5)
θ_{TP}	Take-profit threshold (relative to initial HODL value)	0.10
θ_{SL}	Stop-loss threshold (relative to initial HODL value)	0.20
JIT liquidity provider (MEV searcher)		
N_{JIT}	Target up to top- N swaps per block (by input size)	3
α_{JIT}	Target fraction of active liquidity temporarily added	0.90
Initialization		
$N_{\text{bin}}^{(0)}$	Binomial hill depth used to seed initial liquidity distribution	450
$L_{\text{tot}}^{(0)}$	Total initial liquidity deployed across seed LPs	5.0×10^5
Fee controller		
f_0	Baseline taker fee (static mode / floor for dynamics)	1.0×10^{-4}
$f_{\text{min}}, f_{\text{max}}$	Fee bounds enforced by controller	$(1.0 \times 10^{-5}, 5.0 \times 10^{-2})$
h_{fee}	EWMA half-life (blocks) for controller signals	2
k_{basis}	Toxicity-to-fee multiplier	5.0×10^{-4}
Δf_{min}	Min absolute fee change to commit	5.0×10^{-8}
Δf_{max}	Max fee change per update step	0.01
τ_{cool}	Fee update cooldown (blocks)	0


Table 1: Baseline simulation parameters shared across all models.

4.1 Microstructure diagnostics


Before comparing fee policies, we validate that the ABM reproduces the expected microstructure channel: block-time latency generates transient CEX–DEX deviations that are harvested by priority arbitrage, implying adverse selection (LVR) for LPs.

Figure 1 compares the reference (CEX) price with the AMM (DEX) price and the implied no-arbitrage band. The DEX series exhibits pronounced spikes because the CEX price evolves within each block while on-chain swaps are executed at the block boundary. [More broadly, this channel is consistent with artificial-market evidence that latency and high-frequency priority mechanisms can generate short-horizon volatility and instability even when they improve price alignment in normal periods \[26, 20, 40\].](#) Under static fees, mispricings are then corrected by arbitrage at the beginning of the next block, generating the spike-and-reversion dynamics documented empirically in [34]. Under toxicity-based fees, the same stale-quote gaps are often absorbed by a wider no-arbitrage band, so fewer economically sensible arbitrage swaps are executed and the visible pull of the DEX price back toward the CEX price is attenuated. In the representative Model 0 seed used for the zoom, executed arbitrage swaps fall from 9,725 under static fees to 6,571 under toxicity fees, while in-band arbitrage no-ops rise from 3,275 to 6,428. Sparse or low-liquidity tick regions can also stop an arbitrage swap before its unconstrained target; this liquidity-friction mechanism is present under both fee schedules, whereas the comparative attenuation in panel (b) is driven by the wider dynamic-fee no-arbitrage wedge. The apparent lag between the CEX series and the shaded band is structural: the band is computed from the lagged snapshot m_{t-1} , whereas the plotted m_t is the contemporaneous end-of-block reference price after intra-block diffusion.

Another informative diagnostic is the autocorrelation function (ACF) of end-of-block DEX returns. Figure 2 contrasts the static-fee baseline with the toxicity controller. Under static fees, the ACF shows a pronounced negative spike at lag 1, indicating short-horizon mean reversion across consecutive blocks. A qualitatively similar lag-1 anticorrelation has been documented in [34] and is often linked to arbitrage dynamics; our simulations show that block-time non-atomic CEX–DEX arbitrage alone can generate this signature. Under toxicity-based fees, the lag-1 spike is strongly attenuated in the same representative seed, consistent with the lower rate of economically sensible arbitrage corrections documented in Figure 1.



(a) End-of-block CEX and DEX prices with the no-arbitrage band (representative static-fee run).



(b) Zoomed comparison of the same Model 0 seed and parameters under static fees (left) and toxicity-based fees (right).

Figure 1: Microstructure illustration in Model 0. Panel (a) shows the static-fee CEX and DEX prices with the no-arbitrage band. Panel (b) shows that toxicity-based fees widen the no-arbitrage band during stale-quote episodes, reducing executed arbitrage and attenuating the next-block correction of the DEX price toward the CEX price.




Figure 2: Autocorrelation function of end-of-block DEX log returns in the representative Model 0 seed used in Figure 1. Static fees generate a pronounced lag-1 negative spike (-0.450); toxicity-based fees attenuate the lag-1 autocorrelation to -0.049 , consistent with fewer next-block arbitrage corrections.

4.2 Cross-scenario outcome summary

We now summarize outcomes across all (model, fee schedule) combinations using aggregated cross-scenario figures. This presentation focuses on the paper’s central question: whether dynamic fees can offset block-time adverse selection without rendering the AMM uncompetitive, and how the answer changes once liquidity provision becomes more strategic.

Hedged PnL across models and fee modes. Figure 3 reports the final hedged PnL by cohort (token-1 numéraire). The static-fee benchmark remains unfavorable for standing liquidity in every model. Passive LP hedged PnL is negative in Model 0 (-578.6 ± 4.7), Model 1 (-660.3 ± 3.8), and Model 2 (-666.7 ± 3.5), with active LPs being much more exposed to adverse selection in Model 1 (-5558.1 ± 50.6) and in Model 2 (-5501.7 ± 45.5). This establishes the baseline microstructure result at the aggregate level: a fixed fee does not compensate stale-quote losses, and concentration amplifies that exposure.

Dynamic fees materially change this ranking, but only when the fee signal is sufficiently aligned with adverse selection. In Model 0, all dynamic controllers turn passive LP PnL positive, with the toxicity schedule showing the highest improvement (26698.4 ± 45.3), followed by the DEX-volatility (18356.9 ± 42.9) and CEX-volatility rules (5431.0 ± 7.8). Once active liquidity is introduced (Model 1), the toxicity controller remains the only specification that protects both LP cohorts simultaneously, delivering 1356.0 ± 11.2 for passive LPs and 5704.9 ± 103.4 for active LPs. By contrast, the two volatility-based rules still leave active LPs substantially negative (-3439.6 ± 76.0 and -2658.4 ± 91.3), even though passive LPs become mildly positive.

The same ordering persists in Model 2, but with a sharper incidence pattern. Toxicity-driven fees keep both passive and active LPs positive (496.6 ± 6.4 and 1853.1 ± 66.4), whereas the volatility-based rules keep passive LPs positive but leave active LPs deeply negative. At the same time, the Jitter is essentially inactive under static fees (5.8 ± 0.3) but becomes strongly profitable under every dynamic schedule, especially under toxicity (4086.7 ± 21.9 , versus 2913.3 ± 10.1 and 2138.6 ± 11.8 for the two volatility rules). The updated figure therefore sharpens the paper’s main trade-off: dynamic fees can mitigate LVR for standing liquidity, but they also enlarge the rent pool available to fast, event-contingent liquidity. In this comparison, the toxicity controller is the only rule that robustly

restores positive hedged profitability for both passive and active LPs once strategic liquidity is present.




Figure 3: Final hedged PnL (token-1) across models and fee modes. Cells report mean $\pm$ standard error across 100 seeds. Inactive cohorts (e.g., JIT in Models 0–1) are left blank.

Competitiveness: DEX share. The smart router trades off execution quality against fees, so higher endogenous fees can reduce AMM competitiveness. Table 2 reports three routing-and-fee diagnostics used to evaluate this trade-off: the smart-router DEX share, the cumulative fee value earned by standing LP cohorts, and the time-averaged taker fee. The first diagnostic shows a systematic competitiveness cost. Static fees yield the highest DEX share in every model (35.50% in Model 0, 39.11% in Model 1, and 38.53% in Model 2), whereas dynamic controllers lower routed DEX share to roughly 22.77%–27.94%. The toxicity rule produces the largest decline in routed flow, especially in Model 0, consistent with its higher average fee level. The volatility-based rules preserve slightly more order flow than toxicity, but they still remain well below the static benchmark. Importantly, the AMM remains economically relevant in all configurations; dynamic fees reduce DEX share rather than eliminating routed flow. The improvement in LP protection documented in Figure 3 is therefore achieved through a measurable, but partial, loss of competitiveness rather than through a collapse in trading activity. It is worth noting, however, that the tracked DEX share should be seen as a lower bound for the AMM’s competitiveness, since the smart router is a simplified proxy for real-world routing logic and may not capture all the nuances of how traders choose between venues in practice. For instance, a trader can choose to swap on the DEX for other reasons beyond pure price and fee considerations, such as privacy, censorship resistance, or simply as a part of a bigger strategy that involves on-chain interactions.

Fee levels and fee incidence across cohorts. The remaining columns of Table 2 report the time-averaged taker fee and the corresponding fee-incidence decomposition across passive and active standing LP cohorts. Toxicity is systematically the highest-fee dynamic regime, reaching 14.48 bps in Model 0 and about 11–12 bps in Models 1 and Model 2. The DEX-volatility controller is intermediate, at 7.25 bps in Model 0 and about 4.40–5.22 bps thereafter, while the CEX-volatility controller remains close to 5.3 bps across all models. In Model 0, where all standing liquidity is passive, higher fees map almost one-for-one into passive fee income: cumulative standing-LP fee value rises from 54.7 ± 0.2 under static fees to 532.0 ± 4.1 under toxicity.

Once passive and active liquidity coexist, the fee distribution becomes highly asymmetric. In Models 1 and 2, the active cohort receives the dominant share of dynamic fee income, especially under toxicity, where cumulative standing-LP fee value rises to 1524.1 ± 92.2 in Model 1 and 1303.8 ± 9.2 in Model 2. This comparison is important because it shows that large fee revenues do not, by themselves, guarantee protection against adverse selection. Under both volatility-based rules, active LPs collect substantial fees yet remain strongly negative in hedged PnL terms (Figure 3). What matters is not only charging more, but charging more precisely in the states in which stale-quote risk is highest. The toxicity controller is the only specification that simultaneously raises the fee wedge and moves active LP hedged PnL into positive territory. Once JIT liquidity is admitted, part of the enlarged rent pool is still diverted to the fast liquidity provider, as evidenced by the large positive JIT PnL in Figure 3. The aggregate message is therefore two-sided: dynamic fees can mitigate LVR, but the signal design determines whether the additional fee income compensates standing LPs or is shared with strategic, latency-sensitive liquidity.

Table 2: Routing and fee outcomes across model specifications and fee schedules. The reported diagnostics are the smart-router DEX share, cumulative fee value, and time-averaged taker fee. Entries are means $\pm$ standard errors across 100 independent seeds. DEX share is measured as the percentage of smart-router trades routed to the AMM; passive, active, and total fee values are cumulative token-1 fee revenues for the displayed standing LP cohorts, with total equal to passive plus active fees; average taker fees are reported in basis points. Static rows use the fixed baseline fee $f_0 = 1$ bp.

Model	Fee mode	DEX share (%)	Passive fees	Active fees	Total fees	Avg. fee (bps)
Model 0	Static	35.50 ± 0.30	54.71 ± 0.19	–	54.71 ± 0.19	1.00 (fixed)
Model 0	Toxicity	22.77 ± 0.19	531.99 ± 4.06	–	531.99 ± 4.06	14.48 ± 0.08
Model 0	Volatility (DEX)	23.85 ± 0.42	325.14 ± 11.96	–	325.14 ± 11.96	7.25 ± 0.16
Model 0	Volatility (CEX)	25.05 ± 0.46	247.47 ± 0.59	–	247.47 ± 0.59	5.33 ± 0.02
Model 1	Static	39.11 ± 0.16	19.55 ± 0.03	208.99 ± 2.92	228.53 ± 2.89	1.00 (fixed)
Model 1	Toxicity	26.11 ± 0.57	176.04 ± 8.09	1348.07 ± 84.13	1524.11 ± 92.22	11.72 ± 0.15
Model 1	Volatility (DEX)	27.87 ± 0.67	97.49 ± 0.75	570.81 ± 9.73	668.30 ± 10.48	5.22 ± 0.14
Model 1	Volatility (CEX)	27.25 ± 0.61	94.16 ± 5.23	950.92 ± 10.97	1045.09 ± 16.20	5.36 ± 0.01
Model 2	Static	38.53 ± 0.39	20.63 ± 1.83	196.85 ± 1.12	217.47 ± 0.71	1.00 (fixed)
Model 2	Toxicity	26.45 ± 1.04	103.05 ± 4.70	1200.77 ± 4.53	1303.82 ± 9.23	11.01 ± 0.07
Model 2	Volatility (DEX)	27.94 ± 0.23	66.54 ± 4.11	495.62 ± 1.88	562.16 ± 2.23	4.40 ± 0.12
Model 2	Volatility (CEX)	26.53 ± 0.27	82.41 ± 0.25	854.24 ± 11.71	936.65 ± 11.97	5.31 ± 0.02

Table 3: Arbitrage-related metrics reported as (mean, standard deviation) across simulation runs for the baseline parameterization in Table 1. “Arb Exec” denotes the number of executed arbitrage trades, “No Arb” the number of instances in which arbitrage was infeasible because prices lay within the no-arbitrage band, and “Arb Rej” the number of arbitrage attempts rejected by the profitability filter.

Model	Arb Exec				No Arb				Arb Rej			
	static	vol dex	vol cex	toxicity	static	vol dex	vol cex	toxicity	static	vol dex	vol cex	toxicity
Model 0	9020, 75	5892, 60	6360, 86	5843, 104	3980, 75	7108, 59	6640, 86	7157, 104	0.10, 0.30	0.04, 0.20	0.10, 0.30	0.06, 0.24
Model 1	8696, 62	6008, 65	6360, 86	5529, 97	4304, 62	6992, 65	6640, 86	7471, 97	0.16, 0.36	0.06, 0.23	0.10, 0.30	0.12, 0.32
Model 2	8680, 62	6025, 61	6092, 74	5518, 99	4319, 62	6975, 61	6908, 73	7482, 99	0.08, 0.27	0.06, 0.23	0.04, 0.19	0.10, 0.36

4.3 LVR vs fees vs block size

In this section we discuss the behaviour of the per-block ratio

$$R_t = \frac{\Delta \text{LVR}_t}{\Delta F_t} \quad \text{where} \quad \Delta \text{LVR}_t = \text{LVR}_t - \text{LVR}_{t-1}, \quad \Delta F_t = F_t - F_{t-1} \quad (64)$$

as the size of the block increases. It quantifies how much the fee earned by the LP offset the accumulated LVR in each block. [The LVR theory implies that holding other factors fixed, the important scaling is \[32\]:](#)

$$\mathbb{E}[\text{LVR over an interval of length } \Delta t] \propto \sigma^2 \Delta t. \quad (65)$$

In the simulator, increasing B makes each block represent a longer interval before the end-of-block price m_t is realized and before the pool is compared against it. [This block-size experiment is analogous to latency and market-speed counterfactuals in artificial-market studies of high-frequency trading \[26, 20\].](#) Larger blocks therefore produce larger typical price changes within a block (variance scales with time), which increases the magnitude of mispricing and arbitrage extraction in that block, consistent with (65). This naturally produces approximately linear growth of ΔLVR_t with B . However, in general, even if ΔLVR_t grows with B , the ratio R_t will only remain stable if fee revenue also grows proportionally. Agents’ arrivals are parameterized as per-micro-step rates (for both liquidity takers and LP review/refresh events), so that the expected number of fee-generating trades and LP actions in a block scales approximately linearly with B .

Figures 4 and 5 report the results of this analysis for the four Model 2 fee schedules. Because the plotted object is the median of block-level ratios, values below one should be read as improved marginal fee coverage, not as a guarantee that cumulative hedged PnL over the entire run is positive. The main findings are:

- **Static fee mode:** both passive and active LP cohorts start with median ratios around 2.16–2.20 at $B = 2$ and remain above one for almost the entire sweep, crossing below one only at the extreme right tail ($B \approx 14$ for passive LPs and $B \approx 15$ for active LPs). Fees therefore lag LVR over most of the block-size range, and the active cohort is slightly more exposed than the passive one. However, the ΔLVR per block grows for both passive and active LPs, even if for the former the growth is monotone, while for the latter it is positive between $B > 2$ and $B < 6$ and then starts to decline. This suggests that, as we expected, LVR is mitigated also by just having more uniform flow.
- **Toxicity fee mode:** this is the earliest and cleanest improvement for standing liquidity. The passive and active cohorts are already near break-even at $B = 2$ (1.07 and 1.08, respectively),

fall below one by $B = 3$ –4, and decline further to about 0.54–0.55 at $B = 16$. The toxicity signal therefore raises fees in the same blocks in which stale-price risk is concentrated, rather than only after volatility has accumulated. However, ΔLVR in this case grows monotonically with B for both cohorts, a symptom of the fact that arbitrageurs are less active and the AMM quotes remain outdated for longer. In general we have this behaviour for all dynamic fee schedules, see Figure 6 for an example.

- **Volatility fee modes:** both volatility-based schedules produce a monotone decline in the coverage ratio, but they react more slowly than toxicity. The CEX-volatility rule crosses below one at $B \approx 4$ for passive LPs and $B \approx 5$ for active LPs, while the DEX-volatility rule crosses at about $B \approx 5$ for both cohorts. By $B = 16$ the DEX-volatility controller ends slightly lower (≈ 0.40 –0.42) than the CEX-volatility one (≈ 0.45 –0.47), but both remain less targeted than toxicity at short block times.
- **Jiter:** ratios are reported conditional on successful JIT executions. Under static fees the number of successful JIT opportunities is very small, so the estimated ratio is extremely noisy and the median is strongly negative. Under dynamic fees the Jiter ratio is much more stable and generally stays below one: around $0.57 \rightarrow 0.49$ under toxicity as B moves from 2 to 16, and close to zero at short blocks under the two volatility rules before turning modestly positive at longer blocks. This is consistent with a strategy that captures elevated fees while keeping inventory exposure very short-lived.




Figure 4: Median of $R_t = \Delta LVR_t / \Delta F_t$ vs block time B for the static benchmark and the toxicity-driven fee schedule in Model 2 (50 runs). The Jiter curve is computed conditional on successful JIT executions.




Figure 5: Median of $R_t = \Delta LVR_t / \Delta F_t$ vs block time B for the two volatility-driven fee schedules in Model 2 (50 runs). The Jitter curve is computed conditional on successful JIT executions.




Figure 6: Median per-block ΔLVR vs block time B for active LPs, passive LPs, and the Jiter in Model 2 (50 runs), shown only for the static benchmark and the toxicity-driven fee schedule. These panels make visible the distinction emphasized in the text: static fees combine worse fee coverage with lower or non-monotone ΔLVR dynamics for some cohorts, whereas toxicity improves fee coverage while allowing per-block stale-quote losses to rise more steadily with block length.

5 Conclusions

This paper introduced a granular agent-based model for studying the economics of Uniswap v3 liquidity provision under realistic blockchain frictions and heterogeneous strategic behavior. The model couples a concentrated-liquidity AMM to a reference CEX whose mid-price follows a stochastic-volatility Heston dynamics with permanent impact induced by trades, and embeds a block-based settlement layer with mempool latency and stochastic transaction ordering. Within this environment, arbitrageurs enforce a no-arbitrage band shaped by taker fees and flash-loan costs, while liquidity takers split across an uninformed noise-trader component and a best-execution smart router that endogenizes the DEX share as fees and liquidity conditions change. Liquidity providers are modeled as cash-budgeted agents who mint/burn positions with heterogeneous review clocks, range-selection rules, and risk-management triggers. The framework implements a discrete-time rebalancing benchmark that preserves self-financing and inventory matching, enabling a clean decomposition of outcomes into fees versus Loss-Versus-Rebalancing (LVR) and thus a direct measurement of “hedged PnL” as the economically relevant profitability metric for liquidity provision. [Methodologically, the paper contributes to the artificial financial-market literature by treating AMM fee adaptation as a market-design intervention whose effects are mediated by heterogeneous liquidity takers, arbitrageurs, routers, and liquidity providers \[36, 13, 39, 33\].](#)

Across the baseline configurations, the simulations confirm a central microstructure mechanism: in a block-based system, transient within-block deviations between DEX and CEX prices generate systematic arbitrage opportunities that translate into LVR, especially when liquidity is concentrated near the current price and the AMM effectively posts stale quotes. In Model 0, a static fee produces persistently negative hedged PnL for passive LPs, despite a non-trivial fraction of flow routed to the DEX by the smart router. Allowing the taker fee to adapt according to an underlying signal changes the balance between fee income and adverse selection. Among the tested controllers, the toxicity-based schedule, constructed from an excess-basis proxy measuring how much the DEX price gap exceeds the contemporaneous fee band, post effectively internalizes arbitrage intensity: it raises fees specifically in regimes where stale-price risk is highest, leading to positive hedged PnL for passive LPs while maintaining a stable (slightly reduced) DEX market share. Volatility-based schedules also improve LP outcomes relative to the static benchmark, but generally deliver weaker gains, consistent with volatility being a noisier and less targeted proxy for the realized adverse selection channel driving LVR in this microstructure.

Model 1 highlights a core trade-off created by concentrated liquidity. Introducing active LPs increases depth near the prevailing price, improving DEX competitiveness and raising the router’s DEX share; however, the same concentration amplifies exposure to LVR, pushing active LP hedged PnL below that of passive LPs under most fee regimes. Notably, the toxicity-driven controller is the only specification in which both passive and active LP cohorts achieve positive hedged profits, underscoring that where and when fees respond matters more than simply charging higher fees on average. In other words, adaptive pricing can be beneficial, but only if the control signal is aligned with the microstructural source of adverse selection rather than with broad market variability alone.

Model 2 introduces an MEV searcher performing Just-In-Time (JIT) liquidity. The results emphasize a second, protocol-design-relevant tension: mechanisms that protect LPs from arbitrage (dynamic fees) can simultaneously increase the returns to latency/priority strategies that appropriate fee flow. In the static-fee setting, the JIT agent is approximately break-even, while other LP cohorts remain negative on a hedged basis. Under dynamic fees, the JIT strategy becomes profitable, because it can time liquidity to capture elevated fees while keeping inventory exposure short-lived, yet passive and active LPs fail to attain positive hedged PnL under any of the tested schedules. This outcome is economically coherent: when large swaps are systematically wrapped by JIT liquidity,

the fee stream that would otherwise compensate longer-horizon LPs is diverted to the searcher precisely during the most lucrative blocks, leaving residual LPs to bear LVR without sufficient fee income.

Overall, the experiments support three broad conclusions. First, LVR is fundamentally a microstructure cost, magnified by discrete-time execution and latency, and it is better evaluated through modeling the execution layer and not just via continuous-time diffusion intuition. Second, dynamic fees can improve the sustainability of liquidity provision, but effectiveness depends on using signals that are tightly linked to adverse selection (e.g., excess basis/toxicity) and on managing the endogenous response of order routing. Third, MEV/JIT dynamics can dominate the distribution of fee revenue, potentially negating the benefits of adaptive fees for ordinary LPs unless fee mechanisms and execution rules are co-designed with MEV-resilience in mind.

These findings suggest several directions for future work within the same ABM framework: (i) calibrating agent intensities and parameter regimes to specific pools and chains to connect qualitative insights to empirical magnitudes; (ii) incorporating explicit gas costs, priority fees, and auction-style ordering to model the full MEV supply chain; (iii) extending fee controllers to include anti-JIT features (e.g., time-weighted fee attribution, liquidity “stickiness” requirements, or penalties for ultra-short-lived liquidity) while quantifying welfare impacts on takers and routers; and (iv) studying multi-pool routing and cross-pool arbitrage, where dynamic fees in one venue propagate strategically to others. Taken together, the ABM developed here provides a controlled, interpretable testbed showing that adaptive pricing is a promising—but not standalone—tool for improving LP economics in concentrated-liquidity AMMs operating under realistic blockchain microstructure.

References

- [1] H. Adams, N. Zinsmeister, and D. H. Robinson. Uniswap v2 core, 2020.
- [2] H. Adams, N. Zinsmeister, M. Salem, R. Keefer, and D. Robinson. Uniswap v3 Core Whitepaper. <https://uniswap.org/whitepaper-v3.pdf>, 2021.
- [3] Algebra Finance. Algebra Integral: Adaptive Fee Technical Reference. <https://docs.algebra.finance/algebra-integral-technical-reference/adaptive-fee>, 2026. Accessed 2026-01-09.
- [4] Algebra Protocol. Algebra Protocol Tech Paper. <https://docs.algebra.finance/Algebra%20Protocol%20Tech%20paper.pdf>, 2026. Accessed 2026-01-09.
- [5] G. Angeris, H.-T. Kao, R. Chiang, C. Noyes, and T. Chitra. An analysis of uniswap markets. *arXiv preprint arXiv:1911.03380*, 2020.
- [6] M. Avellaneda and S. Stoikov. High-frequency trading in a limit order book. *Quantitative Finance*, 8(3):217–224, 2008.
- [7] W. A. Brock, C. H. Hommes, and F. O. O. Wagener. More hedging instruments may destabilize markets. *Journal of Economic Dynamics and Control*, 33(11):1912–1928, 2009.
- [8] A. Canidio and R. Fritsch. Arbitrageurs’ profits, lvr, and sandwich attacks: batch trading as an amm design response. *arXiv preprint arXiv:2307.02074*, 2025.
- [9] A. Carro, R. Toral, and M. San Miguel. Markets, herding and response to external information. *PLoS ONE*, 10(7):e0133287, 2015.

- [10] Á. Cartea, F. Drissi, and M. Monga. Decentralised finance and automated market making: Execution and speculation. *SSRN Electronic Journal*, 2022. Working paper.
- [11] Á. Cartea, F. Drissi, and M. Monga. Decentralised finance and automated market making: Predictable loss and optimal liquidity provision. *SSRN Electronic Journal*, 2023. Working paper.
- [12] Á. Cartea, S. Jaimungal, and J. Penalva. *Algorithmic and High-Frequency Trading*. Cambridge University Press, Cambridge, 2015.
- [13] A. Chakraborti, I. M. Toke, M. Patriarca, and F. Abergel. Econophysics review: Ii. agent-based models. *Quantitative Finance*, 11(7):1013–1041, 2011.
- [14] S. N. Cohen, M. Sabate-Vidales, D. Siska, and L. Szpruch. Inefficiency of cfms: Hedging perspective and agent-based simulations, 2023. SSRN Working Paper <https://ssrn.com/abstract=4352120>.
- [15] Curve Finance. CryptoSwap / Curve v2 Whitepaper. <https://curve.fi/files/crypto-pools-paper.pdf>, 2026. Accessed 2026-01-09.
- [16] Curve Finance. Stableswap-NG Pools: Dynamic Fee Formula. <https://docs.curve.finance/stableswap-exchange/stableswap-ng/pools/overview/>, 2026. Accessed 2026-01-09.
- [17] P. Daian, S. Goldfeder, T. Kell, Y. Li, X. Zhao, I. Bentov, L. Breidenbach, and A. Juels. Flash boys 2.0: Frontrunning in decentralized exchanges, miner extractable value, and consensus instability. In *2020 IEEE Symposium on Security and Privacy (SP)*, pages 910–927, 2020.
- [18] Z. Fan, F. Marmolejo-Cossío, D. J. Moroz, M. Neuder, R. Rao, and D. C. Parkes. Strategic liquidity provision in uniswap v3. In *5th Conference on Advances in Financial Technologies (AFT 2023)*, 2023.
- [19] R. Fritsch and A. Canidio. Measuring arbitrage losses and profitability of amm liquidity. In *Companion Proceedings of the ACM Web Conference 2024*, pages 1761–1767, 2024.
- [20] K. Gao, P. Vytelingum, S. Weston, W. Luk, and C. Guo. High-frequency financial market simulation and flash crash scenarios analysis: An agent-based modelling approach. *Journal of Artificial Societies and Social Simulation*, 27(2):8, 2024.
- [21] O. Guéant. *The Financial Mathematics of Market Liquidity: From Optimal Execution to Market Making*. CRC Press, Boca Raton, FL, 2016.
- [22] M. Hafner and H. Dietl. Impermanent loss conditions: An analysis of decentralized exchange platforms. *arXiv preprint arXiv:2401.07689*, 2024.
- [23] L. Heimbach, E. Schertenleib, and R. Wattenhofer. Risks and returns of uniswap v3 liquidity providers. In *Proceedings of the 4th ACM Conference on Advances in Financial Technologies (AFT '22)*, 2022.
- [24] S. L. Heston. A closed-form solution for options with stochastic volatility with applications to bond and currency options. *The Review of Financial Studies*, 6(2):327–343, 1993.
- [25] T. Ho and H. R. Stoll. The dynamics of dealer markets under competition. *The Journal of Finance*, 38(4):1053–1074, 1983.

- [26] S. J. Leal, M. Napoletano, A. Roventini, and G. Fagiolo. Rock around the clock: An agent-based model of low- and high-frequency trading. *Journal of Evolutionary Economics*, 26(1):49–76, 2016.
- [27] LFJ (Trader Joe). Developer Docs: Liquidity Book Fees. <https://developers.lfj.gg/concepts/fees>, 2026. Accessed 2026-01-09.
- [28] B. Maire and M. Wunsch. Market neutral liquidity provision. *Ledger*, 9:115–134, 2024.
- [29] G. Maitrier, G. Loeper, K. Kanazawa, and J.-P. Bouchaud. The ‘double’ square-root law: evidence for the mechanical origin of market impact using tokyo stock exchange data. *Quantitative Finance*, 26(4):491–503, 2026.
- [30] K. Mannaro, M. Marchesi, and A. Setzu. Using an artificial financial market for assessing the impact of tobin-like transaction taxes. *Journal of Economic Behavior & Organization*, 67(2):445–462, 2008.
- [31] Meteora. DLMM Formulas: Fee Calculation Formulas. <https://docs.meteora.ag/overview/products/dlmm/dlmm-formulas>, 2026. Accessed 2026-01-09.
- [32] J. Milionis, C. C. Moallemi, T. Roughgarden, and A. L. Zhang. Automated market making and loss-versus-rebalancing. *arXiv preprint arXiv:2208.06046*, 2023.
- [33] T. Mizuta. An agent-based model for designing a financial market that works well. In *2020 IEEE Symposium Series on Computational Intelligence (SSCI)*, pages 400–406, 2020.
- [34] D. M. D. Nosse, F. Gatta, F. Lillo, and S. Jaimungal. Deviations from tradition: Stylized facts in the era of defi, 2025.
- [35] P. Pellizzari and F. Westerhoff. Some effects of transaction taxes under different microstructures. *Journal of Economic Behavior & Organization*, 72(3):850–863, 2009.
- [36] M. Raberto, S. Cincotti, S. M. Focardi, and M. Marchesi. Agent-based simulation of a financial market. *Physica A: Statistical Mechanics and its Applications*, 299(1–2):319–327, 2001.
- [37] B. Tóth, Y. Lempérière, C. Deremble, J. de Lataillade, J. Kockelkoren, and J.-P. Bouchaud. Anomalous price impact and the critical nature of liquidity in financial markets. *Phys. Rev. X*, 1:021006, Oct 2011.
- [38] X. Wan and A. Adams. Just-in-time liquidity on the uniswap protocol. *Available at SSRN 4382303*, 2022.
- [39] F. Westerhoff and R. Dieci. The effectiveness of keynes–tobin transaction taxes when heterogeneous agents can trade in different markets: A behavioral finance approach. *Journal of Economic Dynamics and Control*, 30(2):293–322, 2006.
- [40] I. Yagi, Y. Masuda, and T. Mizuta. Analysis of the impact of high-frequency trading on artificial market liquidity. *IEEE Transactions on Computational Social Systems*, 7(6):1324–1334, 2020.
- [41] H. Zhang, X. Chen, and L. F. Yang. Adaptive liquidity provision in uniswap v3 with deep reinforcement learning. *arXiv preprint arXiv:2309.10129*, 2023.